

Virtual Filesystem Directory Service Specification

Status of This Memo

This memo provides information to the Grid community about file system directory services. It does not define any standards or technical recommendations. Distribution is unlimited.

Copyright Notice

Copyright © Global Grid Forum (2004). All Rights Reserved.

Abstract

This document describes the specification a file system directory service, which will be one of the essential services for Grid file systems or virtual file systems in the Grid environment. It manages the namespace of federated and virtualized data from file system resources, access control mechanisms, and meta-data management. This document proposes a set of operations needed to be supported by file system directory services. For scalable, large-scale and distributed file system directory management, this document also proposes two types of federation of file system directory services.

Contents

Abstract.....	1
1. Introduction.....	2
2. Virtual Filesystem Directory Services	2
2.1 Virtual Filesystem Directory Service Objects.....	3
2.2 Operations of the Virtual Filesystem Directory Service	5
2.3 Description of WSDL Complex Types	11
2.4 Operation Sequence of VFDS	21
3. Federation of Virtual Filesystem Directory Services.....	21
3.1 Service Referrals.....	22
3.2 Delegated Resolution	22
4. Considerations.....	22
5. Summary and conclusion	23
6. Appendix: Sample WSDL for VFDS	23
Author Information.....	26
Intellectual Property Statement.....	26
Full Copyright Notice	26
References.....	27

1. Introduction

Data in the Grid can be of any format and be stored in any type of storage system. There can be many hundreds of petabytes of data in grids, among which a very large percentage is stored in files. A standard mechanism to describe and organize file-based data is essential for facilitating access to this large amount of data. The Grid File System Working Group (GFS-WG) was established in GGF data area to standardize a mechanism to address this need by providing a Grid File System (GFS) or virtual file system in the Grid environment.

Two major deliverables of the WG are (1) architecture of Grid File System Services and (2) specification of File System Directory Services. File system directory services will manage the namespace of federated and virtualized data from file system resources, access control mechanisms, and meta-data management [1]. It will provide features such as (a) virtualized hierarchical namespaces for files or potentially other types of data (such as live data feeds), (b) efficient and transparent file sharing, and (c) ability to describe and manage file-system and application-specific metadata.

This document intends to describe the design of the file system directory service. It proposes a set of operations needed to be supported by file system directory services. For scalable, large-scale and distributed file system directory management, it also proposes two types of federation of the file system directory services.

The overall architecture of the Grid File System will be specified later in GFS-WG, which provides infrastructure of virtual file systems facilitating federation and sharing of virtualized data from file systems in the Grid environment by using file system directory services.

2. Virtual Filesystem Directory Services

The Virtual Filesystem Directory Service, which will henceforth be referred to as VFDS, enables construction of a uniform, global, hierarchical namespace.[1] This directory service enables federation of individual files as well as filesystem trees that are exported by a variety of distributed file systems. It is comprised of two fundamental namespace components: *virtual directories* and *junctions*. These two essential namespace components, also referred to as *directory entries*, are employed to federate existing files and filesystem trees and construct a uniform hierarchy.

A virtual directory is a VFDS entry that is represented as a filesystem directory, however does not have any corresponding position in any physical filesystem; hence it is *virtual*. A virtual directory, therefore, is purely a namespace entity that functions in much the same way as a conventional filesystem directory by maintaining a list of subentries, which thereby demonstrate a hierarchical relationship.

A junction is a VFDS entry that interconnects an existing file or filesystem tree into the global namespace. It functions in much the same way as traditional distributed file system mount points, with the unique property of maintaining uniform namespace representation while facilitating logical as well as physical referencing to external resources. Junctions are therefore categorized into three basic types: *physical junctions*, *logical junctions*, and *referrals*. A physical junction maintains a URL pointer to the physical location of a filesystem tree or single file; an example URL might look something like: `gsiftp://gridftp.abc.com/filesetX`. In contrast, a logical junction contains a URL that points to a secondary service, like a Replica Location Service (RLS), for logical-to-physical resolution given a logical name; an example URL might look something like:

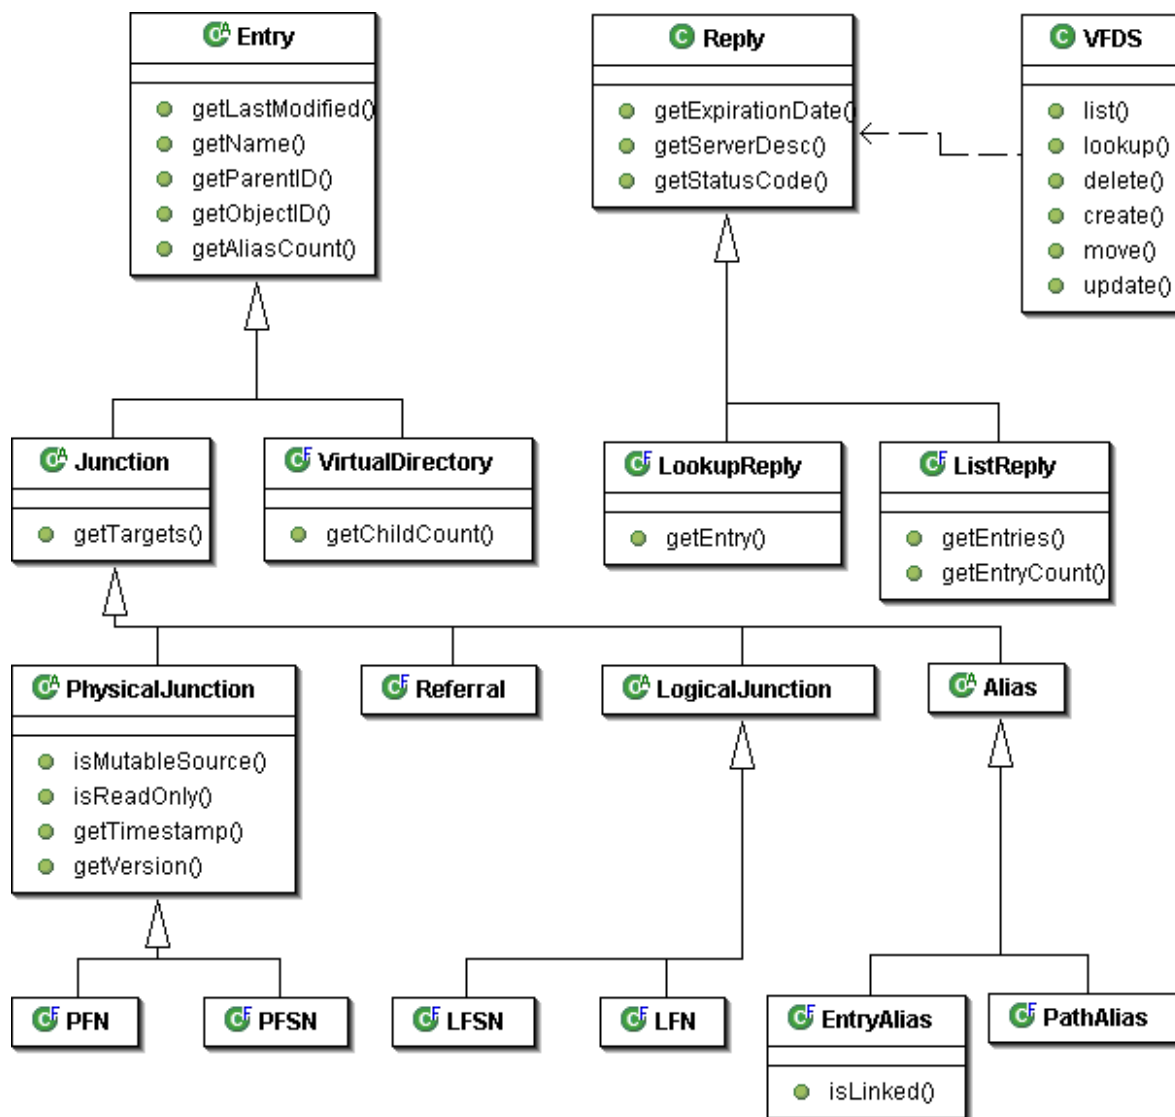
rls://rls.abc.com/logical_name_xyz. Referrals are junctions that link to other VFDS instances, thereby facilitating such features as federation of independent domains of control, scalability of a single domain of control, availability of redundant service instances that may or may not be geographically distributed, etc; an example referral URL might look something like: vfds://vfds.abc.com/. In all cases, junctions are capable of maintaining a list of references (URL pointers) per entry, that is a single junction may render several available URLs, each of which represent replicas or copies of the same data or service instance.

The following sections explore the objects and interface definitions that exemplify the operations of VFDS. This material is not comprehensive, is subject to change, and does not examine the internal procedures of the interface. Error and status codes are not specified at this time.

2.1 Virtual Filesystem Directory Service Objects

VFDS employs an object oriented service interface that consists of two chief objects: *Reply* and *Entry*. The service port type corresponds to the principal service class, that is the *VFDS* class, which defines the interface methods or port type definition in WSDL. These methods package all reply messages in a *Reply* objects. Therefore, all interface methods return a *Reply* object to the requester in response to both mutator and accessor operations. The *Reply* object encapsulates the following basic members: *status/error code*, *time-to-live expiration*, *array of Entry objects* (if applicable), and *textual server description*. The *Reply* object therefore not only discloses information consequent from the request made, but also delivers all *Entry* objects corresponding to the response if the request was made to an accessor method. The *Entry* object simply represents the VFDS entry maintained by the directory service.

Following is an object diagram of the VFDS objects represented in this document:



As illustrated here, and previously mentioned, an *Entry* object can be represented as either a virtual directory or a junction. This is shown in the second row of objects in the object diagram as a subclass relationship. Notice that the *Junction* class, however, is an abstract class and is further extended by its subclasses, which represent the various types of junctions previously mentioned. Finally, notice that both the physical and logical classes are abstract and also extended to be represented by concrete classes that denote file or filesystem junctions, which are physical or logical respectively. Incidentally, these classes are as follows: PFN=Physical File Name, PFSN=Physical File System Name, LFSN=Logical File System Name, and LFN=Logical File Name.

2.2 Operations of the Virtual Filesystem Directory Service

The basic operations of VFDS are:

- 1) Lookup operation, governed by access control permissions, to convert a globally unique path name to the corresponding VFDS entry.
- 2) Creation, removal, and rename operations for all entry types.
- 3) Operations for managing attributes or status of an entry.

As revealed by the VFDS object diagram, the VFDS service definition specifically contains the following operations:

Operation: create

Description:

Creates a new VFDS entry in the directory service database that corresponds to the entry object provided and positions the new entry at the absolute path specified.

This method is designed to process all defined types of entry objects, so the entry parameter should reflect the appropriate runtime cast of the abstract Entry class.

If the optional *oid* parameter is used, then the *path* argument is treated as a relative path; relative to the VFDS entry denoted by the *oid* parameter.

(similar to a `mkdir` and `mount` in a Unix shell).

Parameters:

path: Absolute path, or relative path if *oid* is used, for the object to create.

entry: A valid Entry object populated with all of the member values necessary to represent the entry to create.

oid: [Optional] Object ID of the Current Working Directory

Returns:

Reply object that contains the status of the requested operation

Signature:

```
public Reply create( String path, Entry entry, [String oid] ) throws
    ProtocolException;
```

WSDL:

```
<operation name='create' parameterOrder='path entry'>
  <input message='tns:VFDS_create__Request_Soap' />
  <output message='tns:VFDS_create_Response_Soap' />
  <fault name='ProtocolException1' message='tns:ProtocolException1' />
</operation>
```

Operation: update

Description:

Updates an existing VFDS entry in the directory service database that corresponds to the entry object provided and overwrites the new entry at the absolute path, or OID, specified.

This method is designed to process all defined types of entry objects, so the entry parameter should reflect the appropriate runtime cast of the abstract Entry class.

If the optional *oid* parameter is used, then the *path* argument is treated as a relative path; relative to the VFDS entry denoted by the *oid* parameter.

Parameters:

path: Absolute path, or relative path if *oid* is used, for the object to update.

entry: A valid Entry object populated with all of the member values necessary to represent the entry to create.

oid: [Optional] Object ID of the Current Working Directory

Returns:

Reply object that contains the status of the requested operation

Signature:

```
public Reply update( String path, Entry entry, [String oid] ) throws
    ProtocolException;
```

WSDL:

```
<operation name='update' parameterOrder='path entry'>
  <input message='tns:VFDS_update__Request_Soap' />
  <output message='tns:VFDS_update_Response_Soap' />
  <fault name='ProtocolException1' message='tns:ProtocolException1' />
</operation>
```

Operation: delete

Description:

Delete the VFDS entry specified by path. If path is null an exception is thrown.

If path represents a virtual directory, all sub-directories must be deleted before it can be deleted.

If the optional *oid* parameter is used, then the *path* argument is treated as a relative path; relative to the VFDS entry denoted by the *oid* parameter.

(similar to a `rm` in a Unix shell).

Parameters:

path: Absolute path, or relative path if *oid* is used, for the object to delete.

oid: [Optional] Object ID of the Current Working Directory

Returns:

Reply object that contains the status of the requested operation

Signature:

```
public Reply delete( String path, [String oid] ) throws  
    ProtocolException;
```

WSDL:

```
<operation name='delete' parameterOrder='path'>  
  <input message='tns:VFDS_delete__Request_Soap' />  
  <output message='tns:VFDS_delete_Response_Soap' />  
  <fault name='ProtocolException' message='tns:ProtocolException' />  
</operation>
```

Operation: <code>list</code>

Description:

Returns a Reply object that represents the directory listing of the specified path that represents a virtual directory; if path is null then an exception is thrown.

If path represents a virtual directory, its contents will be returned; otherwise an exception is thrown.

If the optional *oid* parameter is used, then the *path* argument is treated as a relative path; relative to the VFDS entry denoted by the *oid* parameter.

(similar to a `ls -al` in a Unix shell).

Parameters:

path: The full path, or relative path if *oid* is used, to the virtual directory to be listed.

oid: [Optional] Object ID of the Current Working Directory

Returns:

Reply object containing directory listing (array of Entry objects).

Signature:

```
public Reply list( String path, [String oid] ) throws
    ProtocolException;
```

WSDL:

```
<operation name='list' parameterOrder='path'>
  <input message='tns:VFDS_list__Request_Soap' />
  <output message='tns:VFDS_list_Response_Soap' />
  <fault name='ProtocolException3' message='tns:ProtocolException3' />
</operation>
```


Operation: lookup

Description:

Returns a Reply object that contains an Entry object, which represents the VFDS entry and associated attributes of the specified path; if path is null then an exception is thrown.

If the path denotes a virtual directory, then only the Entry object that represents that VFDS virtual directory is returned; not its sub-entries. (see list() for a listing of sub-entries)

If the optional *oid* parameter is used, then the *path* argument is treated as a relative path; relative to the VFDS entry denoted by the *oid* parameter.

(similar to a `ls file` or `stat()`).

Parameters:

path: The full path, or relative path if *oid* is used, to the entry whose attributes are desired.

oid: [Optional] Object ID of the Current Working Directory

Returns:

Reply object containing the Entry object requested.

Signature:

```
public Reply lookup( String path, [String oid] ) throws
    ProtocolException;
```

WSDL:

```
<operation name='lookup' parameterOrder='path'>
  <input message='tns:VFDS_lookup__Request_Soap' />
  <output message='tns:VFDS_lookup_Response_Soap' />
  <fault name='ProtocolException4' message='tns:ProtocolException4' />
</operation>
```

Operation: move

Description:

Moves or renames the VFDS entry denoted by src to dest.

Both parameters are required.

(similar to a `mv` in a Unix shell)

Parameters:

src: Absolute path to source or original VFDS entry.

dest: Absolute path to destination for the VFDS entry denoted by src.

Returns:

Reply object that contains the status of the requested operation

Signature:

```
public Reply move( String src, String dest, [String srcOID], [String  
destOID] ) throws ProtocolException;
```

WSDL:

```
<operation name='move' parameterOrder='src dest'>  
  <input message='tns:VFDS_move__Request_Soap' />  
  <output message='tns:VFDS_move_Response_Soap' />  
  <fault name='ProtocolException2' message='tns:ProtocolException2' />  
</operation>
```

2.3 Description of WSDL Complex Types

The objects used as components of the communication interface defined by WSDL, also referred to as *complex types*, are the two chief service interface objects previously itemized: *Reply* and *Entry*. These two objects serve as integral components in the task of communicating and representing the VFDS service.

2.3.1 Reply Object

The following describes the interface definition for the VFDS *Reply* object:

Method: <code>getExpirationDate</code>
--

Description:

Returns the expiration timestamp of this Reply, that is the time at which the Reply should be considered stale. This property is analogous to time-to-live (TTL).

Returns:

The expiration date of this Reply, if available and appropriate; otherwise a null value is used to indicate this reply is a point-in-time reply

Signature:

```
public Date getExpirationDate();
```

WSDL:

```
<xsd:complexType name="Reply">
  ...
  <xsd:sequence>
    <xsd:element name="expirationDate" type="xsd:dateTime" nillable="true"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: <code>getStatusCode</code>

Description:

Returns the status or error code; otherwise referred to as a reply code. Every reply from the VFDS service will contain some status or error code and is readily available via this method. A value of zero '0' indicates success, while a value other than zero may represent an error.

Returns:

Status/error code

Signature:

```
public int getStatusCode();
```

WSDL:

```
<xsd:complexType name="Reply">
  ...
  <xsd:sequence>
    <xsd:element name="statusCode" type="xsd:int"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: <code>getServerDesc</code>

Description:

Returns a simple character string that represents the text that VFDS server implementations use to identify themselves to VFDS service users for usage tracking and other purposes.

Returns:

VFDS server implementation description

Signature:

```
public String getServerDesc();
```

WSDL:

```
<xsd:complexType name="Reply">
  ...
  <xsd:sequence>
    <xsd:element name="serverDesc" type="xsd:string" nillable="true"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

2.3.1.1 Reply Subclass: LookupReply

The VFDS *LookupReply* class is a concrete class that is a direct subclass of *Reply*. It appends one accessor method to the structural definition of *Reply* for the purpose of enabling access to enclosed content:

Method: <code>getEntry</code>

Description:

Returns the Entry object representing the result of the call made. This method will return null if there were no Entry objects returned by the remote call made.

Returns:

The Entry objects contained within this Reply object

Signature:

```
public Entry getEntry();
```

WSDL:

```
<xsd:complexType name="Reply">
  ...
  <xsd:sequence>
    <xsd:element name="entry" type="tns:Entry" nillable="true"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

2.3.1.2 *Reply Subclass: ListReply*

The VFDS *ListReply* class is a concrete class that is a direct subclass of *Reply*. It appends two accessor methods to the structural definition of *Reply* for the purpose of enabling access to enclosed content:

Method: <code>getEntryCount</code>

Description:

Returns the number of entries contained within this Reply object.

Returns:

The number of entries contained within this Reply object

Signature:

```
public int getEntryCount();
```

WSDL:

```
<xsd:complexType name="Reply">
  ...
  <xsd:sequence>
    <xsd:element name="entryCount" type="xsd:int"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: <code>getEntries</code>

Description:

Returns an array of Entry objects representing the result of the call made. This method will return null if there were no Entry objects returned by the remote call made.

Returns:

The array of Entry objects contained within this Reply object

Signature:

```
public Entry[] getEntries();
```

WSDL:

```
<xsd:complexType name="Reply">
  ...
  <xsd:sequence>
    <xsd:element name="entries" type="tns:ArrayOfEntry" nillable="true"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

2.3.2 *Entry Object*

All valid VFDS *Entry* objects are characterized by the following abstract class definition, which is to say that the base *Entry* class is abstract and serves as a structural definition for its various subclass heirs.

The following describes the interface definition for the VFDS *Entry* object:

Method:	<code>getLastModified</code>
---------	------------------------------

Description:

Returns the last modified timestamp of the VFDS entry represented by this object. This timestamp reflects the last time the corresponding VFDS entry was modified within the namespace repository and does not reflect the last modified timestamp from any physical target.

Returns:

Modification timestamp of the VFDS entry this object represents

Signature:

```
public Date getLastModified();
```

WSDL:

```
<xsd:complexType name="Entry">
  ...
  <xsd:sequence>
    <xsd:element name="lastModified" type="xsd:dateTime" nillable="true"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method:	<code>getName</code>
---------	----------------------

Description:

Returns the name of the VFDS entry this object represents.

Returns:

Name of the represented VFDS entry

Signature:

```
public String getName();
```

WSDL:

```
<xsd:complexType name="Entry">
  ...
  <xsd:sequence>
    <xsd:element name="name" type="xsd:string" nillable="true"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: getObjectID

Description:

Returns the session specific object ID that uniquely identifies the VFDS entry this object represents.

Returns:

Object ID of the represented VFDS entry

Signature:

```
public String getObjectID();
```

WSDL:

```
<xsd:complexType name="Entry">
  ...
  <xsd:sequence>
    <xsd:element name="oid" type="xsd:string" nillable="false"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: getParentID

Description:

Returns the session specific object ID of the parent object that uniquely identifies the parent VFDS entry of the VFDS entry this object represents.

Returns:

Parent object ID of the represented VFDS entry

Signature:

```
public String getParentID();
```

WSDL:

```
<xsd:complexType name="Entry">
  ...
  <xsd:sequence>
    <xsd:element name="pid" type="xsd:string" nillable="false"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: <code>getAliasCount</code>

Description:

Returns the number of aliases this VFDS entry has associated with it. An alias is a VFDS entry that symbolizes another VFDS entry via a different pathname.

Returns:

Number of the associated aliases

Signature:

```
public int getAliasCount();
```

WSDL:

```
<xsd:complexType name="Entry">
  ...
  <xsd:sequence>
    <xsd:element name="aliasCount" type="xsd:ArrayOfInt" nillable="false"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

2.3.2.1 Entry Subclass: VirtualDirectory

The VFDS *VirtualDirectory* class is a concrete class that is a direct subclass of *Entry*. It appends one method to the structural definition of *Entry*:

Method: <code>getChildCount</code>

Description:

Returns the number of immediate subentries related to this VirtualDirectory in the directory hierarchy.

Returns:

Number of immediate subentries related to this VirtualDirectory

Signature:

```
public int getChildCount();
```

WSDL:

```
<xsd:complexType name="VirtualDirectory">
  ...
  <xsd:sequence>
    <xsd:element name="childCount" type="xsd:int"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```


2.3.2.2 *Entry Subclass: Junction*

The VFDS *Junction* class is an abstract class that is a direct subclass of *Entry*. It appends one method to the structural definition of *Entry*.

Method:	<code>getTargets</code>
---------	-------------------------

Description:

Returns at least one String, typically a URL, that identifies the target resource for the VFDS junction represented by this object. The URL syntax must comply with RFC 1738.

Returns:

An array of URLs or end point references that refer to target resources denoted by this junction

Signature:

```
public String[] getTargets();
```

WSDL:

```
<xsd:complexType name="Junction">
  ...
  <xsd:sequence>
    <xsd:element name="targetURL" type="xsd:ArrayOfString" nillable="false"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

2.3.2.2.1 *Junction Subclass: Referral*

The VFDS *Referral* class is an abstract class that is a direct subclass of *Junction*. A *Referral* is a VFDS entry that refers to another VFDS service instance. This enables delegate VFDS service instances with the potential of separation of administrative ownership and control.

2.3.2.2.2 *Junction Subclass: Alias*

The VFDS *Alias* class is an abstract class that is a direct subclass of *Junction*. An *Alias* is a VFDS entry that symbolizes another VFDS entry. It is effectively an “alias” or alternate name for the entry being symbolized, and may have a different path and/or name that uniquely identifies it in the global namespace.

2.3.2.2.2.1 *Junction Subclass: EntryAlias*

The VFDS *EntryAlias* class is a concrete class that is a direct subclass of *Alias*. An *EntryAlias* persistently maintains its referential relationship with the VFDS entry it is aliasing by object ID. Since entry aliases point to target entries by their respective object IDs, the targeted VFDS entry may be moved and/or renamed without disintegrating the alias entry.

Entry aliases truly symbolize the VFDS entry that they target by simply presenting an alternate name to the same entry. Provided that the user creating the alias entry has the appropriate authorization, the establishment of the alias entry will transactionally increment a link count attribute of the targeted entry. If the user creating the alias entry does not have adequate permission to alter

the link count attribute of the target entry, then the alias is created without a corresponding “link count” reference.

Method: <code>isLinked</code>

Description:

Identifies whether or not the targeted VFDS entry object has a corresponding link count associated with this *EntryAlias* object.

Returns:

True if the target VFDS entry’s link count attribute was incremented; false if it was not.

Signature:

```
public boolean isLinked();
```

WSDL:

```
<xsd:complexType name="EntryAlias">
  ...
  <xsd:sequence>
    <xsd:element name="linked" type="xsd:boolean"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

2.3.2.2.2 Junction Subclass: *PathAlias*

The VFDS *PathAlias* class is a concrete class that is a direct subclass of *Alias*. A *PathAlias* refers to a target VFDS entry by pathname. If the VFDS entry object that represents the target pathname is moved, renamed, or deleted, the *PathAlias* object reference will become invalid.

2.3.2.2.3 Junction Subclass: *LogicalJunction*

The VFDS *LogicalJunction* class is an abstract class that is a direct subclass of *Junction*. A *LogicalJunction* is a VFDS entry that junctions a resource by its unique logical name, which implies the necessity of a Logical-to-Physical resolver in order to obtain a physical location of the data represented; an example resolver might be something like a *Replica Location Service*.

2.3.2.2.4 Junction Subclass: *PhysicalJunction*

The VFDS *PhysicalJunction* class is an abstract class that is a direct subclass of *Junction*. It appends four methods to the structural definition of *Junction*:

Method:	<code>isMutableSource</code>
---------	------------------------------

Description:

Identifies whether or not the file or filesystem source targeted by this VFDS entry can change.

Returns:

True if the data source is mutable; false if it is not.

Signature:

```
public boolean isMutableSource();
```

WSDL:

```
<xsd:complexType name="PhysicalJunction">
  ...
  <xsd:sequence>
    <xsd:element name="mutableSource" type="xsd:boolean"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method:	<code>isReadOnly</code>
---------	-------------------------

Description:

Identifies whether or not a local copy of the data targeted by this VFDS entry should be locally read-only.

Returns:

True if the local copy of the data should be read-only; false if the local copy allows writes.

Signature:

```
public boolean isReadOnly();
```

WSDL:

```
<xsd:complexType name="PhysicalJunction">
  ...
  <xsd:sequence>
    <xsd:element name="readOnly" type="xsd:boolean"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: getTimestamp

Description:

This represents a replica or fileset's point-in-time timestamp corresponding to the time at which the source snapshot was made.

Returns:

The point-in-time timestamp of the targeted data

Signature:

```
public Date getTimestamp();
```

WSDL:

```
<xsd:complexType name="PhysicalJunction">
  ...
  <xsd:sequence>
    <xsd:element name="timestamp" type="xsd:dateTime" nillable="true"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

Method: getVersion

Description:

This represents a version number of the targeted data if available.

Returns:

Version number of the targeted data if available

Signature:

```
public float getVersion();
```

WSDL:

```
<xsd:complexType name="PhysicalJunction">
  ...
  <xsd:sequence>
    <xsd:element name="version" type="xsd:float"/>
    ...
  </xsd:sequence>
</xsd:complexType>
```

2.4 Operation Sequence of VFDS

This section will briefly outline the operational sequence of communicating and interacting with VFDS.

2.4.1 List entries of a virtual directory

To retrieve a list of directory entries for a given virtual directory in the global namespace hierarchy, the *list* operation is utilized. The client making such a request would simply perform a *list* service operation against the VFDS service, whose returning messages will be constructed into a *Reply* object. The requesting client then retrieves the list of all subentries, via the *getEntries* method, immediately related to the virtual directory specified by the path used in the *list* operation. At this point the requesting client can process any number (or all) of the returned *Entry* objects. If an unexpected exception occurs, the status code available via the *Reply* object can be checked.

In summary, the following operations are necessary to retrieve a directory listing in VFDS:

- 1) Invoke the *list* operation of the VFDS service, specifying the globally unique path of the virtual directory to list.
- 2) Retrieve all *Entry* objects from the *Reply* object by calling the *getEntries* method.
- 3) Enumerate the *Entry* objects delivered via the *Reply* object.

2.4.2 Create a virtual directory

To create a virtual directory in the global namespace hierarchy, the *create* operation is utilized. The client placing such a request would simply perform a *create* service operation against the VFDS service, specifying the global path and providing an instantiated *VirtualDirectory* object as the *Entry* parameter. Once the service completes this operation, its returning messages will be constructed into a *Reply* object. The requesting client then evaluates the status code available via the *getStatusCode* method of the *Reply* object. If an unexpected exception occurs, the status code available via the *Reply* object can be checked.

In summary, the following operations are necessary to retrieve a directory listing in VFDS:

- 1) Construct a *VirtualDirectory* object with the appropriate property values assigned (name, modification time, etc.).
- 2) Invoke the *create* operation of the VFDS service, specifying the globally unique path of the virtual directory to create and providing the previously instantiated *VirtualDirectory* object from step 1 as the *Entry* parameter.
- 3) Evaluate the status code returned by the service via the *getStatusCode* method of the *Reply* object.

3. Federation of Virtual Filesystem Directory Services

A global namespace service directly implies the employment of a multitude of namespace servers by virtue of geographical distribution, segregated domains of ownership and control, scalability, and redundancy/availability. A principal goal of a global namespace service is to provide a location independent view of consistent access paths to data. Since these access paths are represented by hierarchal path names, symbolizing a globally unique identifier of data, it is a natural extension of the design to postulate an architecture that federates multiple namespace servers in a hierarchical fashion. Similar to the well established DNS model, VFDS servers can be interlinked by referrals whilst providing a seamless and transparent view of the namespace. Once several instances of the namespace service are interlinked, the most obvious challenge is related to path name resolution

when dealing with paths that cross referral boundaries. There are two fundamental approaches to resolving path names that span multiple namespace domains or service instances: service referrals and delegated resolution.

3.1 Service Referrals

The most straightforward and arguably the most secure and truly scalable approach to resolving path names that span multiple domains or service instances is to place the onus of handling VFDS referrals on the VFDS client. In this approach, the namespace server would simply return a VFDS referral to the VFDS client when a junction to another namespace server is encountered. The client implementing the VFDS API is then responsible for continuing the task of resolving the original path name by connecting to the namespace server indicated by the VFDS referral and querying the newly connected server for further (relative) path name resolution.

One clear advantage of this approach is the direct management of namespace service connections, which implies authentication and authorization control per connection, rather than accessing a referred namespace server via proxied security. Additionally, this approach promotes distributed work load balancing; instead of requiring VFDS servers to handle namespace requests for both locally managed namespace and remotely managed namespace via proxy.

3.2 Delegated Resolution

Another possible approach to resolving path names that span multiple domains or service instances is to empower the VFDS server to delegate queries to other VFDS servers for complete resolution of any given path. Although this approach is demonstrated in DNS, it should be noted that the security requirements are quite different. Since DNS generally operates in a public read-only manner without authentication and authorization per DNS server, it is not too unreasonable to endorse such an approach. VFDS, however, facilitates the possibility of requiring authentication per service instance and enforcing access control per entry. Nevertheless, an approach that allows for the possibility of delegated resolution should be considered as at least an optional mode of operation; incidentally DNS is capable of both approaches.

4. Considerations

There are several issues to consider, with respect to VFDS, that have not been explored in this document.

- ? Security – The topic of security as a whole is not discussed in this specification document. Security is recognized as a substantial area of interest and will require further investigation.
- ? Replication of VFDS databases – To enhance fault tolerance and reliability, replication of namespace service data is indispensable. The consistency model required by VFDS needs to be investigated.
- ? Backup – Backup of VFDS data may be required.
- ? Discussion of access control lists (ACLs) within VFDS, their purpose, scope, representation, and enforcement. If access permissions defined by physical filesystems are to be represented within VFDS then significant consideration must be taken with respect to consistency problems between access permissions of a virtual file and the corresponding file data.

- ? Removal or modification of a file data without notification to the file system directory services.
- ? Consistency problems between file data replicas.
- ? Interoperability issue with NFSv4 and CIFS.

5. Summary and conclusion

This document intended to describe the specification of the Virtual Filesystem Directory Service, which will be one of the essential services for the realization of a Grid File System. It manages the namespace of federated and virtualized data from file system resources, access control mechanisms, and meta-data management.

This document proposed a set of operations needed to be supported by VFDS. Additionally, it proposed two approaches to federation of VFDS service instances for scalable, large-scale and distributed namespace management.

Further detailed discussion for specification and evaluation by implementing file system directory services are needed with respect to security, performance, consistency, scalability, and reliability. The evaluation needs to consider functionality of a client library, especially, with and without client attribute cache.

6. Appendix: Sample WSDL for VFDS

The following is a sample Web Services Description Language (WSDL) file intended to describe VFDS in terms of a web service using a standard XML format.

```
<service name='VFDS'>
  <port name='VFDS' binding='tns:VFDS'>
    <soap:address location='http://vfds.almaden.ibm.com:6060/VFDS/' />
  </port>
</service>
<portType name='VFDS'>
  <operation name='create' parameterOrder='path entry'>
    <input message='tns:VFDS_create__Request_Soap' />
    <output message='tns:VFDS_create_Response_Soap' />
    <fault name='ProtocolException1' message='tns:ProtocolException1' />
  </operation>
  <operation name='delete' parameterOrder='path'>
    <input message='tns:VFDS_delete__Request_Soap' />
    <output message='tns:VFDS_delete_Response_Soap' />
    <fault name='ProtocolException' message='tns:ProtocolException' />
  </operation>
  <operation name='list' parameterOrder='path'>
    <input message='tns:VFDS_list__Request_Soap' />
    <output message='tns:VFDS_list_Response_Soap' />
    <fault name='ProtocolException3' message='tns:ProtocolException3' />
  </operation>
  <operation name='lookup' parameterOrder='path'>
    <input message='tns:VFDS_lookup__Request_Soap' />
    <output message='tns:VFDS_lookup_Response_Soap' />
    <fault name='ProtocolException4' message='tns:ProtocolException4' />
  </operation>
  <operation name='move' parameterOrder='src dest'>
    <input message='tns:VFDS_move__Request_Soap' />
    <output message='tns:VFDS_move_Response_Soap' />
  </operation>
</portType>
</service>
```

```

        <fault name='ProtocolException2' message='tns:ProtocolException2' />
    </operation>
</portType>
<binding name='VFDS' type='tns:VFDS'>
    <soap:binding transport='http://schemas.xmlsoap.org/soap/http' style='document' />
    <operation name='create'>
        <map:java-operation name='create' signature='xxx'>
            <map:fault name='ProtocolException1' java-
type='com.ibm.vfds.ProtocolException' />
        </map:java-operation>
    <soap:operation
soapAction='http://systinet.com/wsdl/com/ibm/vfds/VFDS#create?xxx' style='document' />
    <input>
        <soap:body parts='path entry' use='literal' />
    </input>
    <output>
        <soap:body parts='response' use='literal' />
    </output>
    <fault name='ProtocolException1'>
        <soap:fault name='ProtocolException1' use='literal' />
    </fault>
</operation>
    <operation name='delete'>
        <map:java-operation name='delete' signature='xxx'>
            <map:fault name='ProtocolException' java-
type='com.ibm.vfds.ProtocolException' />
        </map:java-operation>
    <soap:operation
soapAction='http://vfds.almaden.ibm.com/wsdl/com/ibm/vfds/VFDS#delete?xxx'
style='document' />
    <input>
        <soap:body parts='path' use='literal' />
    </input>
    <output>
        <soap:body parts='response' use='literal' />
    </output>
    <fault name='ProtocolException'>
        <soap:fault name='ProtocolException' use='literal' />
    </fault>
</operation>
    <operation name='list'>
        <map:java-operation name='list' signature='xxx'>
            <map:fault name='ProtocolException3' java-
type='com.ibm.vfds.ProtocolException' />
        </map:java-operation>
    <soap:operation
soapAction='http://vfds.almaden.ibm.com/wsdl/com/ibm/vfds/VFDS#list?xxx' style='document' />
    <input>
        <soap:body parts='path' use='literal' />
    </input>
    <output>
        <soap:body parts='response' use='literal' />
    </output>
    <fault name='ProtocolException3'>
        <soap:fault name='ProtocolException3' use='literal' />
    </fault>
</operation>
    <operation name='lookup'>
        <map:java-operation name='lookup' signature='xxx'>
            <map:fault name='ProtocolException4' java-
type='com.ibm.vfds.ProtocolException' />
        </map:java-operation>
    <soap:operation
soapAction='http://vfds.almaden.ibm.com/wsdl/com/ibm/vfds/VFDS#lookup?xxx'
style='document' />
    <input>
        <soap:body parts='path' use='literal' />

```



```
</input>
<output>
  <soap:body parts='response' use='literal' />
</output>
<fault name='ProtocolException4'>
  <soap:fault name='ProtocolException4' use='literal' />
</fault>
</operation>
<operation name='move'>
  <map:java-operation name='move' signature='xxx'>
    <map:fault name='ProtocolException2' java-
type='com.ibm.vfds.ProtocolException' />
  </map:java-operation>
  <soap:operation
soapAction='http://vfds.almaden.ibm.com/wsdl/com/ibm/vfds/VFDS#move?xxx' style='document' />
  <input>
    <soap:body parts='src dest' use='literal' />
  </input>
  <output>
    <soap:body parts='response' use='literal' />
  </output>
  <fault name='ProtocolException2'>
    <soap:fault name='ProtocolException2' use='literal' />
  </fault>
</operation>
</binding>
```

Author Information

Osamu Tatebe
Grid Technology Research Center, AIST
1-1-1 Umezono, Tsukuba
Ibaraki 3058568 Japan
o.tatebe@aist.go.jp

Manuel Pereira, Leo Luan, Ted Anderson
IBM Almaden Research Center
650 Harry Road
San Jose, CA 95120, USA
mpereira@us.ibm.com
leoluan@us.ibm.com
ota@us.ibm.com

Jane Xu
IBM Systems and Technology Group
5600 Cottle Road
San Jose, CA 95193, USA
jxu@us.ibm.com

Intellectual Property Statement

The GGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the GGF Secretariat.

The GGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the GGF Executive Director.

Full Copyright Notice

Copyright (C) Global Grid Forum (2004). All Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included on all such copies and derivative works. However, this document itself may not be modified in any way, such as by removing the copyright notice or references to the GGF or other organizations, except as needed for the purpose of developing Grid Recommendations in which case the procedures for copyrights defined in the GGF Document process must be followed, or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by the GGF or its successors or assigns.

This document and the information contained herein is provided on an "AS IS" basis and THE GLOBAL GRID FORUM DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE."

References

- [1] Leo Luan and Ted Anderson, "Grid Namespace for Files", GGF working draft, GGF8, 2003
https://forge.gridforum.org/projects/gfs-wg/document/Grid_Namespace_for_Files/en/1
- [2] S. Shepler, et al., "Network File System (NFS) version 4 Protocol", RFC3530, 2003