

XtreemOS

*Enabling Linux
for the Grid*



Extensions to JSDL

XtreemOS Consortium

Matej Artac (XLAB), Massimo Coppola (CNR), Toni Cortes (BSC),
Yvon Jegou (INRIA), John Mehnert-Spahn (UDUS), Ramon Nou (BSC)

March 2010

XtreemOS IP project

is funded by the European Commission under contract IST-FP6-033576



Information Society
Technologies



- **XtreemOS is a Grid Operating System**
- **Provides transparent access to the Grid**
 - “run on the grid as if you where on your desktop”
- **To execute a job:**
 - Execute the application JSDL on a desktop
- **XtreemOS project supported by the EU**
 - **`www.xtreemos.eu`**

Outline

- **Dynamic behavior in resource selection**
- **Resource selection tolerance**
- **Location of files**
- **Job execution time**
- **Firewalls**
- **Interactive jobs**
- **Credential delegation**
- **Checkpointing**

Managing dynamic behavior

- **XtreemOS allows *physical resource sharing***
 - Multiple **processes, containers** or **VMs** : akin to Clouds
 - Good application matching is important

App A	App B	App A + APP B
<ul style="list-style-type: none">• 30 % CPU load• 50 % RAM load• 20 %Net BW	<ul style="list-style-type: none">• 60% CPU load• 30% RAM load• 30% Net BW	<ul style="list-style-type: none">• 90% CPU load• 80% RAM load• 50% Net BW

- **XtreemOS allows interactive applications**
 - Dynamic app behaviour, dynamic app start/stop
 - Appropriate policies and resource selection strategies
- **Associate a dynamic meaning to most JSDL tags**
 - E.g. Total RAM \leftrightarrow Free RAM
 - “Static” value may differ from “dynamic” value

- **Constraint on **dynamic** resource characteristics**
 - Applications may accept some tolerance
 - needs 4GB of memory, but would run with 3GB free
 - needs a 3Ghz Xeon, but would accept a 30% loaded machine
- **Proposed changes**
 - Resource tags are extended with a new attribute to indicate tolerance
 - Real change is in the RangeValue_Type tag
 - Backward compatible
 - Dynamic threshold value = static bound * tolerance
 - For those attributes whose dynamic semantics is meaningful

- **Comparison with dynamic values can be slacker**

```
<jSDL-srds:IndividualDiskSpace>
```

```
<jSDL-srds:Range>
```

```
<jSDL-srds:LowerBound tolerance="0.1">100000000
```

```
</jSDL-srds:LowerBound>
```

```
<jSDL-srds:UpperBound tolerance="0.9">2400000000
```

```
</jSDL-srds:UpperBound>
```

```
</jSDL-srds:Range>
```

```
</jSDL-srds:IndividualDiskSpace>
```

- **Select machines whose**

- disk space amount X is : $100\text{MB} < X < 2400\text{MB}$
- **free** disk space Y is : $0.1 * 100\text{MB} < Y < 2400\text{MB} * 0.9$

Example

```
<jSDL-srds:IndividualDiskSpace>  
  <jSDL-srds:Exact tolerance="0.9">10000000  
  </jSDL-srds:Exact>  
</jSDL-srds:IndividualDiskSpace>
```

- Exact requirements turned into tolerance ranges
- Select machines whose
 - Whose disk space amount X is : $X == 10\text{MB}$
 - Whose *free* disk space Y is : $0.9 * 10\text{MB} < Y < 10\text{MB}$

- **Exclusive access to resources not always needed**
 - Application may be willing to share a CPU
 - Historically reliable machines may be preferred
- **Proposed changes**
 - Add tag `Uptime` (machine uptime in min)
 - Add tag `IdlePercentage` (% of idle cycles)
 - Already dynamic, no tolerance attribute needed
 - Can be combined with other tags possibly using dynamic tolerance
 - Allow to select unloaded, long standing machines

- **Ordinary JSDL use**

```
<jSDL-srds:IdlePercentage><jSDL:Range>  
  <jSDL:LowerBound>0</jSDL:LowerBound>  
  <jSDL:UpperBound>50</jSDL:UpperBound>  
</jSDL:Range></jSDL-srds:IdlePercentage>  
<jSDL-srds:Uptime>  
  <jSDL:LowerBoundedRange>2</jSDL:LowerBoundedRange>  
</jSDL-srds:Uptime>
```

- **Select machines with less than 50% CPU idle**
 - Load is always measured dynamically
- **Refuse machines who have just boot up (>2 min)**

Working implementation

- **XtreemOS employs described extensions**
- **Syntax refined after first prototypes**
- **XML schemata defined extending standard JSDL**
 - Extensions are backward compatible
- **XML validation enforced**

Extension Schema fragment

```
<?xml version="1.0" encoding="UTF-8"?>
<xsd:schema xmlns:xsd="http://www.w3.org/2001/XMLSchema"
  targetNamespace="http://schemas.ggf.org/jsdl/2005/11/jsdl-srds"
  elementFormDefault="qualified"
  xmlns="http://schemas.ggf.org/jsdl/2009/11/jsdl-srds"
  xmlns:jsdl-srds="http://schemas.ggf.org/jsdl/2005/11/jsdl-srds"
  xmlns:jsdl="http://schemas.ggf.org/jsdl/2005/11/jsdl">
  <!-- Import normative schema -->
  <xsd:import namespace="http://schemas.ggf.org/jsdl/2005/11/jsdl"
    schemaLocation="Jsdl_Normative_OGF.xsd"/>
  <!-- COMPLEX TYPES: Definitions for the RangeValueType -->
  <xsd:complexType name="Boundary_Type"
    xmlns="http://schemas.ggf.org/jsdl/2005/11/jsdl-srds">
    <xsd:simpleContent>
      <xsd:extension base="xsd:double">
        <xsd:attribute name="exclusiveBound" type="xsd:boolean" use="optional"/>
        <xsd:attribute name="tolerance" type="xsd:double"
          xmlns="http://schemas.ggf.org/jsdl/2005/11/jsdl-srds" use="required"/>
        <xsd:anyAttribute namespace="##other" processContents="lax"/>
      </xsd:extension>
    </xsd:simpleContent>
  </xsd:complexType>
```

- **Try to get resources close to data**
 - Applications need a way to specify the important files
 - Can be used by scheduler as a hint
 - Not the same as stage in/out files
 - It is complementary
 - Not all files may be important
 - Depending on file system, some stage in/out may be needed
- **Proposed changes**
 - New tag (`SchedulingHint`) with file information

Example

```
<SchedulingHint>  
  <FileSystem type="XtreemFS">  
    <Volume id="Default">  
      <File>out.txt</File>  
      <File>tmpdir/process.txt</File>  
    </Volume>  
  </FileSystem>  
</SchedulingHint>
```


Job time

- **User has little control on when a job is executed**
 - Cannot decide
 - Starting time
 - Days and/or times when a job can be executed
 - Duration (it it takes longer checkpoint and stop)
- **Proposed changes**
 - New tag (LifeTime) constrains on job execution

```
<LifeTime>
  <StartTime> datetime </StartTime>
  <ExecutionTime> time </ExecutionTime>
  <Constraints>
    <Constraint>
      <DayOfWeek> dayofweek </DayOfWeek>
      <TimeInterval>
        <Start> time </Start>
        <End> time </End>
      </TimeInterval>
    </Constraint>
  </Constraints>
</LifeTime>
```

Firewall

- **Firewalls make jobs life nearly impossible**
 - We should offer a way to specify network needs
 - Resource selections should take these needs into account
- **Proposed changes**
 - Add some parameters to the resource tag

Example

```
<jsd1:Resources>  
  <network:Network  
    xmlns:net=http://xtreemos.org/schemas/jsdl/net/2008  
05>  
      <net:Netmask>201.123.123.0/24</net:Netmask>  
      <net:Ports>1000-2000,60000,65000</net:Ports>  
      <net:Proto>TCP</net:Proto>  
    </network:Network>  
</jsdl:Resources>
```

Interactive jobs

- **Interactive jobs should be executable in a Grid**
 - We need to detect which jobs are interactive
 - We need to set the environment for their interactivity
- **Proposed changes**
 - Extend application element to describe interactivity

Example

- Solution 1: extend existing tags

<Input>/dev/tty</Input>

<Output>stdout</Output>

- Add new tags

<X11>true/false</X11>

- Tags need to be analyzed on both client and resource sides

- **XtreemOS jobs can execute Grid requests**
 - Interactive jobs
 - Needs credentials from the user (certificates, ...)
- **Various delegation schemes defined**
 - Proxy certificates
 - XtreemOS SSO service
 - Dtokens
- **When delegation is not handled in user space**
 - Need to be requested from JSDL
- **Need a new tag**

- **Checkpointing is a common case in the Grid**
 - Applications should be able to specify its parameters
 - Who should initiate it (i.e user/system)
 - Which checkpointer to use (i.e. BLRC version X.Y)
 - Which container (i.e. cgroups)
 - Application information (i.e. #procs, #threads, ...)
 - Type of checkpointing protocol (i.e. coordinated or not)
- **Proposed changes**
 - Add a new tag (`JobFaultTolerance`) with all needed parameters

Example

```
<JobFaultTolerance>  
  <Initiator>  
    <User> yes </User>  
    <OperatingSystem> yes </OperatingSystem>  
    <Application>no</Application>  
  </Initiator>  
  
  <Checkpoint> BLCR </Checkpoint>  
  
  <CheckpointVersion>0.8.2  
</CheckpointVersion>  
  ...
```

Example

...

```
<ContainerType>cgroups</ContainerType>
```

```
<ContainerParameter>
```

```
  <NetworkNS>yes</NetworkNS>
```

```
  <PidNS>yes</PidNS>
```

```
  <IpcNS>yes</IpcNS>
```

```
  <Uts>yes</Uts>
```

```
</ContainerParameter>
```

...

Example

...

```
<ApplicationSoftwareResources>  
  <Singleproc>yes</Singleproc>  
  
  <Singlethread>yes</Singlethread>  
  
  <Sysipcshm>yes</Sysipcshm>  
  
  <Sysipcmsgq>yes<Sysipcmsgq>  
  <Sysipcsem>yes<Sysipcsem>  
  <Files>yes<Files>  
</ApplicationSoftwareResources>
```

...

Example

...

```
<CheckpointProtocol>Coordinated</CheckpointProtocol>  
  <CheckpointProtocolParameter>  
    <Periodinseconds>5</Periodinseconds>  
  </CheckpointProtocolParameter>  
</JobFaultTolerance>
```