

Open Cloud Computing Interface - Use cases and requirements for a Cloud API

Status of This Document

This document provides information to the Cloud and Grid community. Distribution is unlimited.

Copyright Notice

Copyright (c) Open Grid Forum (2009). All Rights Reserved.

Abstract

This document is an informal document describing the use cases and requirements for a Cloud API. This document is created as part of the deliverables of the Open Cloud Computing Interface working group. The document is the first deliverable and will be used to demonstrate and validate the features of the Open Cloud Computing Interface.

Contents

TOC comes here

1. Introduction

This document includes all the use cases and requirements which were gathered in the working group. The next two sections hold the use cases and the requirements. Each use case is defined by name and a set of functional and non-functional requirements. All the requirements are categorized and formatted in tables. Each requirement is (whenever possible) mapped to a use case. Priorities have been added to all requirements.

2. OCCI Use Cases

The following section describes the use cases which were gathered during the requirements analyses for the OCCI working group. They are used to set up the requirements and later on to verify the OCCI specification.

2.1. SLA-aware cloud infrastructure using SLA@SOI

There is a need for a standard interface to dynamic infrastructure provisioning. While doing so it must be guaranteed and verified that the infrastructure provisioning uses 'machine-readable' SLAs. (SLA_SOI)

Functional Requirements

- VM Description: request format important - In this area is where there is least coherency amongst providers.
- VM Description: A means to add non-functional constraints on functional attributes.
- VM Management: all parameters in the request should be "monitor-able" and verifiable. Full control of resources (VMs) allocated required; at a minimum: start, stop, suspend, resume.
- VM Monitoring: Monitoring non-functional constraints declared in provisioning request
- Network Management: resources assignable by network tag - defaults of public and private further sub-categorisation could be allowed e.g. tag of web could be assigned to the public network group.
- Storage Management: simple mount points, reuse storage SaaS offerings

Non-functional Requirements

- Security: Transport and user level (ACLs? OAuth?) security

- Quality of Service: Can be many - Part of service offering from the infrastructure provider e.g. Security, QoS, geo-location, isolation levels - NFPs are the basic building blocks of differentiating IaaS providers.
- Scheduling Information: When a particular resource is to be run. Also in which order should a collection of resources be ran in the case that one resource is dependent on another.

2.2. Service Manager to control the Life cycle of Services

This use case is based in the 'Service Manager' (SM) layer of the RESERVOIR project architecture. 'Service Providers' (SP) willing to deploy their service on the Cloud use this layer to control the service life cycle. The SM operates over the Cloud infrastructure automatically as the service demands. In a way, the SM maps the service configuration and needs to calls to the Cloud infrastructure, so many of the requirements imposed by the SM are due to the flexibility that the SM aims to provide to SPs. (RV)

Functional Requirements

- Network Management: There should be methods for the Allocation of private networks, where VMs can be attached to. A special network (e.g. 'Public Network') should be available. When some network interface is attached to it, the infrastructure must assign it a public IP address.
- Image Management: There should be methods to register, upload, update and download disk images.
- VM Description: It should be possible to describe all the VM hardware components and their attributes, along with any restriction regarding the VM location:
 - Memory: Size
 - CPU: Architecture, amount of CPU's, speed.
 - Disk: Size, Interface (SCSI, IDE, SATA...), RAID (yes/no, and RAID level). Disk image to mount. Automatic backup (yes/no, backups frequency...).
 - Network: Interfaces, for each interface its bandwidth, and Network they are attached to.
 - Geographical restrictions: Location(s) where the VM can/cannot be deployed (for example for legal purposes).
 - Migration allowed (yes/no): If migration is supported by the infrastructure, this flag sets if it is allowed for the VM.
- VM Management: There should be methods to allow the SM to change the VM state (for example, from ACTIVE to SUSPENDED), if such transition is allowed by the infrastructure (i.e. is defined in the OCCI's State Machine). The description of a VM can be changed when the machine is running (ACTIVE, SUSPENDED...). But it will not be taken into account until the machine is stopped and started again, unless it is a change regarding geographical or migration restrictions. Each disk backup will have an id, as the images defined by the SM. Methods to download any backup should be provided. As each backup is, after all, a disk image, it should be possible to mount it on any VM. For example, it should be possible to stop a VM, change its configuration so its disk mounts this backup image, and restart the VM.
- Monitoring: The status (We use the term 'status' when talking about monitoring, and try not to use the term 'state' to avoid confusion with the states of the OCCI State Machine.) representation of any element is given as a list of keys and their values. For example, the status of a memory component could be given by the amount of memory used and the cache memory. Then, the keys could be 'used' and 'cache', and their values '142MB'

and '430MB'. Both the request and the reply use the corresponding element identifier. Two types of monitoring should be supported:

- Pull based: The SM can request the status of any element it has registered: VMs, networks... Also, the SM can request the status of components, for example, the status of certain disk of a certain VM.
- Publish/subscribe based: The SM can subscribe to be notified about events on the VMs and/or Networks. Some of the events to be notified are:
 - Errors on some component of a VM.
 - Changes on the state of a VM (e.g. from ACTIVE to SUSPENDED).
 - Periodic notifications about some element state. The frequency of this notifications can be configured in the subscription message.
- Error messages: If some VM could not be created, or some image could not be uploaded, etc... the platform should return an error message carrying a detailed description of the reason.
- Identifications: Networks, VMs and images should have unique IDs, (UUIDs, URIs, or the like). It is to be determined whether components of VMs (disks, memory...) should have an unique ID too. IDs are assigned by the Cloud infrastructure when the corresponding element is created.

Non-functional Requirements

- Both for hardware configuration and monitoring values there should be a clear, standard way to set which magnitude the value represents. For example, when setting the memory size to '2', it must be clear that we refer to GBs and not to MBs. An option would be setting the value to '2GB', another would be allowing to set both the value and the magnitude: value '2' and magnitude 'GB'.
- Protocols: The transport, message format, and state representation should use open and standard protocols, each one which strong software support (i.e. libraries and frameworks available for several programming languages).

2.3. Interoperability across Cloud Infrastructures using OpenNebula

OpenNebula is a Virtual Infrastructure Engine that allows the management of Virtual Machines on a pool of physical resources and/or cloud providers. It offers three main functionalities: backend of a public cloud, tool to manage a virtual infrastructure in the data-center or cluster (private cloud), and tool to achieve cloud interoperation (hybrid cloud), the latter being relevant in this use case. The aim of this use case is to state the requirements that an API for cloud providers should take into account in order to expose an interface that will enable the management of groups of Virtual Machines across them. These requirements are gathered from the experience using OpenNebula as a tool to manage Virtual Machines from different cloud providers. Currently, there are two set of plugins for OpenNebula to access Amazon EC2 and ElasticHosts cloud providers that leverage the use of both cloud providers in a transparent fashion for the end user. (ONE)

Functional Requirements

- VM Description: Virtual Machines should be described consistently across cloud providers using a slim set of indispensable attributes, such as:
 - Memory: Amount of RAM needed by the Virtual Machine
 - CPU: Number of CPUs needed by the Virtual Machine (this needs to be normalized)

- Disk: Disks that will conform the basic filesystem and possibly others for the Virtual Machine
- Network: How many network interface this Virtual Machine should have, and where should be attached
- VM Management: API should offer functionality to enforce operations upon Virtual Machines, such as:
 - DEPLOY: Launches the Virtual Machine
 - SHUTDOWN: Shutdown the Virtual Machine
 - CANCEL: Cancels the Virtual Machine in case of failure, or destroys it if it is running
 - CHECKPOINT: Creates a snapshot of the Virtual Machine
 - SAVE: Creates a snapshot of the Virtual Machine AND suspends it
 - RESTORE: Resumes a Virtual Machine from a previous snapshot
 - POLL: Retrieves information about Virtual Machine state and consumption attributes (percentage of Memory, CPU used, bytes transferred, and so on)
- Additionally, Virtual Machines should be in one of the following states:
 - PENDING: VM is waiting for a physical resource slot.
 - BOOTING: VM is being booted
 - RUNNING: VM is active, it should be able to start offering a service
 - SUSPENDED: VM is suspended, waiting for a resume.
 - SHUTDOWN: VM is being shutdown.
 - CANCEL: VM has been canceled by the user or by a scheduler.
 - FAILED: VM crashed or hasn't started properly.
- Network Management: API should expose functionality to
 - Create Private Virtual Networks
 - Attach Public IP to Virtual Machine
- Image Management: The ability to upload disk images is fundamental to virtual machine management to avoid the need to reinstall software for each cloud provider. The upload process should return an identifier to be used in the Virtual Machine Description.

Non-functional Requirements

- Security: Security should be handled using X509 certificates for authentication. Also, authorization can be based on said certificates and ACL lists.
- Quality of Service: When used in conjunction with Haizea, OpenNebula provides advanced reservation functionality. Cloud providers API should provide similar capabilities to ensure proper QoS.

2.4. AJAX web front-end directly calling API

A cloud provider implements their customer web front-end as an entirely client-side AJAX application calling the OCCI API directly.

Functional Requirements

- Completeness: API must contain complete set of calls to completely specify and control cloud (but this is likely only ~15-20 verbs on ~3-4 nouns!)
- Responsiveness: Calls must return swiftly. In particular, we should provide a simple and quick call to poll the `_list_` of servers, drives, etc. that exist without listing all of their properties, since this is computationally much cheaper for the cloud to return, and will need to be regularly polled to catch any servers, etc. that are created outside of the interface.

Non-functional Requirements

- Syntax: A simple JSON syntax for the API will make the AJAX interface much simpler to implement

2.5. Single technical integration to support multiple service providers

(This is the major benefit ElasticHosts sees from a standardized IaaS API). Today, each cloud provider (ElasticHosts, GoGrid, Amazon, etc.) integrates independently with every other player in the cloud ecosystem (CohesiveFT, RightScale, etc), producing $O(n^2)$ separate technical integrations. In the future, if all cloud providers and cloud ecosystem partners use a single standard API, then we have $O(n)$ technical integrations, and all potential partnerships can immediately interoperate.

Non-functional Requirements

- Uptake: Standardized IaaS API needs strong uptake in by both cloud providers and cloud ecosystem.

2.6. Wrapping EC2 in OCCI

Amazon EC2 is the de-facto standard for IaaS today. While Amazon is unlikely to join the OCCI process at first, we must ensure that the OCCI API would be capable of interfacing to EC2 so that gateways can be written and so that Amazon could join at a later date.

Functional Requirements

- Semantics: Must include the ability to fully describe core EC2 objects and operations

Non-functional Requirements

- Wish list: I'd love to see an open source OCCI <-> EC2 gateway

2.7. Automated Business Continuity and Disaster Recovery

Maintain up-to-date remote shadows of physical and/or virtual machines, such that in the event of a disaster it is possible to start and switch to the remote machines.

Functional Requirements

- VM Description: Metadata mapping to legacy systems
- VM Management: Automated management in the event of a disaster (e.g. startup, IP changes).
- Network Management: Runtime alteration of IPs
- Image Management: Advanced, rsync style updates to synchronise machines with physical equivalents (e.g. rsync block devices to remote raw disk files).

Non-functional Requirements

- Quality of Service: Reservation of capacity sufficient for fail over

2.8. Simple scripting of cloud from Unix shell

An end user wishes to script a simple task (such as starting a server at midnight every night and shutting it down an hour later, automating fail over, reporting, etc.). They are using a typical Unix/Linux setup, so would like to write a simple cron job which carries this out.

Non-functional Requirements

- Syntax: This should be as simple as possible to place minimal barriers to entry on the user. The user should not need any development tools or libraries. They should be able to write 1-2 lines of shell script, posting a simple <5 lines of command data using curl, wget, etc.

2.9. Typical web hosting cluster

An end-user runs a typical web hosting cluster on a cloud, with: 2 database servers, 2 front-end web server (bursting to 3 under load), a load balancer (either a specialized virtual machine or provided by the cloud like GoGrid)

Functional Requirements

- Completeness: The API should be able to fully express this cluster, which will require at least: 5 virtual machines, storage for each virtual machine, two networks (a private one connecting the machines, and the public Internet also connected to the load balancer), a fixed static IP for the website on the public Internet, possible specification of the load balancer itself.

2.10. Manage cloud resources from a centralized dashboard

An end user wishes to view and control all of his cloud-based resources in a lightweight (perhaps AJAX-based) console, perhaps the same web front-end referred to in this use case: AJAX web front-end directly calling API

Functional Requirements

- Completeness: Every resource provided by the cloud is discoverable by the API, and every action that can be performed on all these resources is also available via the API, together with actuators to actually perform those actions, and all the attributes of the resources are available via the API.
- Responsiveness: Calls must return swiftly. In particular, we should provide a simple and quick call to poll the `_list_` of servers, drives, etc. that exist without listing all of their properties, since this is computationally much cheaper for the cloud to return, and will need to be regularly polled to catch any servers, etc. that are created outside of the interface. (text copied from AJAX web front-end directly calling API)
- Categorizability: (there's gotta be a better word...) The client must be able to identify what type each resource is in order to display like-typed resources together and in order to provide separate UI views that might be specialized for certain resource types. For example, the client must be able to differentiate between a compute resource that does not represent an actual CPU (perhaps this is a compute template) and between a compute resource that actually represents a running CPU. The interface for actually-running CPUs might display the current IP address of the instance and allow you to SSH into the instance, while a different tab in the interface might display all the compute templates and allow you to instantiate instances from them.

Non-functional Requirements

- Usability: This should be a user interface with context-menus and context-aware links that allow the user to easily see what actions can be performed for each resource.

2.11. Compute Cloud

A cloud provider implements a RESTful API for provisioning, executing, and monitoring of tasks.

Functional Requirements

- **Secure:** API must be secured to ensure that only authorized identities are permitted to use the API.
- **Resource:** An endpoint must be created for external monitoring, status, and auditing of the task. This endpoint would be responsive to RESTful calls supporting AJAX and other clients.
- **Scripted:** The target system needs to understand and process directives which would be provided with the task. These directives would include the ability to pull binaries or data onto the system, run executables, and status the system resources.

Non-functional Requirements

- **Single Compute Method:** The resultant service should be the same service that can be used for many other purposes. It could be used for monitoring of system health, system life-cycle management, system patching, and configuration changes. If this was the only service on the system initially, it could then be used to build up the other services in a plug-in manner.

3. OCCI Requirements

3.1. Functional Requirements

This section deals with the functional requirements. The requirements have been split up in tables and prioritized.

Table 1. Functional requirements on VM description

ID	Description	Usecases	Priority
A.1.1	Attributes to define memory, CPU, disk and network requirements.	2.2, 2.3, 2.6	High
A.1.2.	Attributes to define placement constraints, such as geographical location	2.2	Medium
A.1.3.	Attributes to allow migrations if is supported by the infrastructure.	2.2	Medium
A.1.4.	The API should be able to fully express a cluster: 5 VMs, storage for each VM, two networks (a private one connecting the machines, and the public internet also connected to the load balancer), a fixed static IP for the website on the public internet, possible specification of the load balancer itself	2.9	High
A.1.5.	A means to add constraints (non-functional, functional) on attributes declared in a provisioning request	2.1	High
A.1.6.	Resource execution scheduling. Allow provisioned resources to execute sometime in the future from the original request	2.1	Medium
A.1.7.	Common operating systems should be supported	-	High
A.1.8.	Resources should be grouped according to provider policies	-	Medium
A.1.9.	On requesting of new resource(s) - the request should be fully complete/describing	-	High

Table 2. Functional requirements on VM management

ID	Description	Usecases	Priority
A.2.1.	Methods to start, stop, suspend and resume VMs	2.1, 2.2, 2.3, 2.5, 2.11, 2.10	High
A.2.2.	Automated management in the event of a disaster	2.1, 2.7	Medium
A.2.3.	Provide ID for each backup disk and for the images	2.2	High
A.2.4.	Provide methods to donwload any backup	2.2	Medium
A.2.5.	API should offer functionality to enforce the following operations: deploy, shutdown, cancel, checkpoint, save, restore, poll (could be merged with monitoring)	2.3	High
A.2.6.	State model should include: pending, booting, running, suspended, shutdown, cancel, failed	2.3	Medium
A.2.7.	Listing collections should complete quickly without listing all properties for each entry	2.4	Medium
A.2.8.	Allow for resource representations be updated and have those changes executed upon VMs	-	Low
A.2.9.	Support the usage of terminal, web, desktop and automated management interfaces	2.10	Low
A.2.10.	Support the migration of resources from a physical resource to the cloud, from a cloud to another cloud and from a virtual resource to the cloud	-	Medium

Table 3. Functional requirements on Network management

ID	Description	Usecases	Priority
A.3.1.	VPN creation	2.3	Low
A.3.2.	Multiple network connection (Public and Private)	2.1, 2.2, 2.3	High
A.3.3.	Runtime IP change/attachment	2.3, 2.7	Medium
A.3.4.	Tagged Multiple network connections	2.1, 2.2, 2.3	Low
A.3.5.	Support network setups which support an 'Intercloud' setup.	-	Medium

Table 4. Functional requirements on Storage management

ID	Description	Usecases	Priority
A.4.1.	Use of URIs as mount points - allows reuse of Storage service offerings	2.1	High
A.4.2.	Allow attachment of additional storage at runtime	-	Medium

Table 5. Functional requirements on Image management

ID	Description	Usecases	Priority
A.5.1.	Methods capable to register, upload, update and download disk images	2.2	Medium
A.5.2.	Updates based on rsync commands to synchronize machines with physical equivalents	2.7	Medium
A.5.3.	When an upload completes successfully, an identifier to it should be returned	2.2	Low

Table 6. Identifications/References

ID	Description	Usecases	Priority
A.6.1.	Unique IDs for VM images and their components	2.2	Medium

Table 7. Monitoring

ID	Description	Usecases	Priority
A.7.1.	Pull-based monitoring that request the status of the elements such as network , VM ...	2.1, 2.2	High
A.7.2.	Publish/subscribe monitoring that request events occur in the VM or networks such as Errors on some component, changes in the VM state and other periodic notifications	2.2	Medium
A.7.3.	Attributes that define simple quick call to poll the list of servers, drives, etc	2.4	Low
A.7.4.	Attributes about resource consumption of the VM from the hypervisor (CPU, memory...)	2.1, 2.2	Low
A.7.5.	Management reports should be generated from in some of the following formats XML, PDF	-	Low

3.2. Non-functional Requirements

This section deals with all the non-funtional requirements.

Table 8. Security requirements

ID	Description	Usecases	Priority
B.1.1.	Use of X509 Certificates	2.3	High
B.1.2.	Use of X509 Certificates	B.1, 2.1	High
B.1.3.	Attributes in the VM description to define Security levels	2.1	High
B.1.4.	Transport and user level security	2.1	High
B.1.5.	Allow geographical region be specified and adhered to	B.4	Medium

Table 9. Quality of Service

ID	Description	Usecases	Priority
B.2.1.	Capacity requirements for recovery / failover	2.7	Low
B.2.2.	Attributes in the VM describes to define QoS level	2.1	Medium
B.2.3.	Attributes in the VM describes to define Isolation level	2.1	Medium
B.2.4.	Attributes in advanced reservation functionality	2.3	Low
B.2.5.	Allow VM response time be specified	B.4	Medium

Table 10. Syntax

ID	Description	Usecases	Priority
B.3.1.	No development tools or libraries for the final user, doing it as simple as possible	2.8	Medium
B.3.2.	Simple JSON syntax make Ajax interface much simpler to implement	2.4, 2.10	Medium
B.3.3.	Clear definition of units (MB, GB etc) used in requests	A.2, 2.4	Medium

Table 11. Backup/Disaster recovery

ID	Description	Usecases	Priority
B.4.1.	Support a backup functionality of cloud resources	-	Low
B.4.2.	The interface should reconsider failover, disaster recovery and business continuity plans	-	Medium

4. Conclusions

The previous chapter described to what extent Cloud requirements for an API exist. What's missing up to now is a general solution that fits most of the needs, is simple to implement, highly responsive (throughput), globally applicable (standard well known as well as dynamic connections), secure, highly recognized (commonly known and implemented) and last but not least standardized.

This working group plans to use these use cases and requirements for the creation of an Cloud API.

5. Contributors

The following people have contributed to the requirements gathering and use case documentation.

Table 12. List of contributors

Name	Affiliation
Michael Behrens	R2AD
Richard Davies	ElasticHosts
Andy Edmonds	SLA@SOI - Intel
Sam Johnston	Australian Online Solutions
Thijs Metsch (Editor)	Sun Microsystems - RESERVOIR project
George Reese	O'Reilly
Luis Roderio-Merino	Telefónica I+D
Shlomo Swidler	MyDrifts
Tino Vázquez	DSA-Research at UCM

6. Intellectual Property Statement

The OGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the OGF Secretariat.

The OGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the OGF Executive Director.

7. Disclaimer

This document and the information contained herein is provided on an "As Is" basis and the OGF disclaims all warranties, express or implied, including but not limited to any warranty that the use of the information herein will not infringe any rights or any implied warranties of merchantability or fitness for a particular purpose.

8. Full Copyright Notice

Copyright (C) Open Grid Forum (2009). All Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included on all such copies and derivative works. However, this document itself may not be modified in any way, such as by removing the copyright notice or references to the OGF or other organizations, except as needed for the purpose of developing Grid Recommendations in which case the procedures for copyrights defined in the OGF Document process must be followed, or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by the OGF or its successors or assignees.

9. References

[SLA_SOI] <http://sla-at-soi.eu/>. SLA@SOI project website

[RV] <http://www.reservoir-fp7.eu/>. RESERVOIR project website

[ONE] <http://opennebula.org/>. OpenNebula website