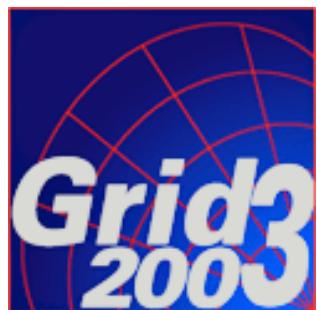


Grid3

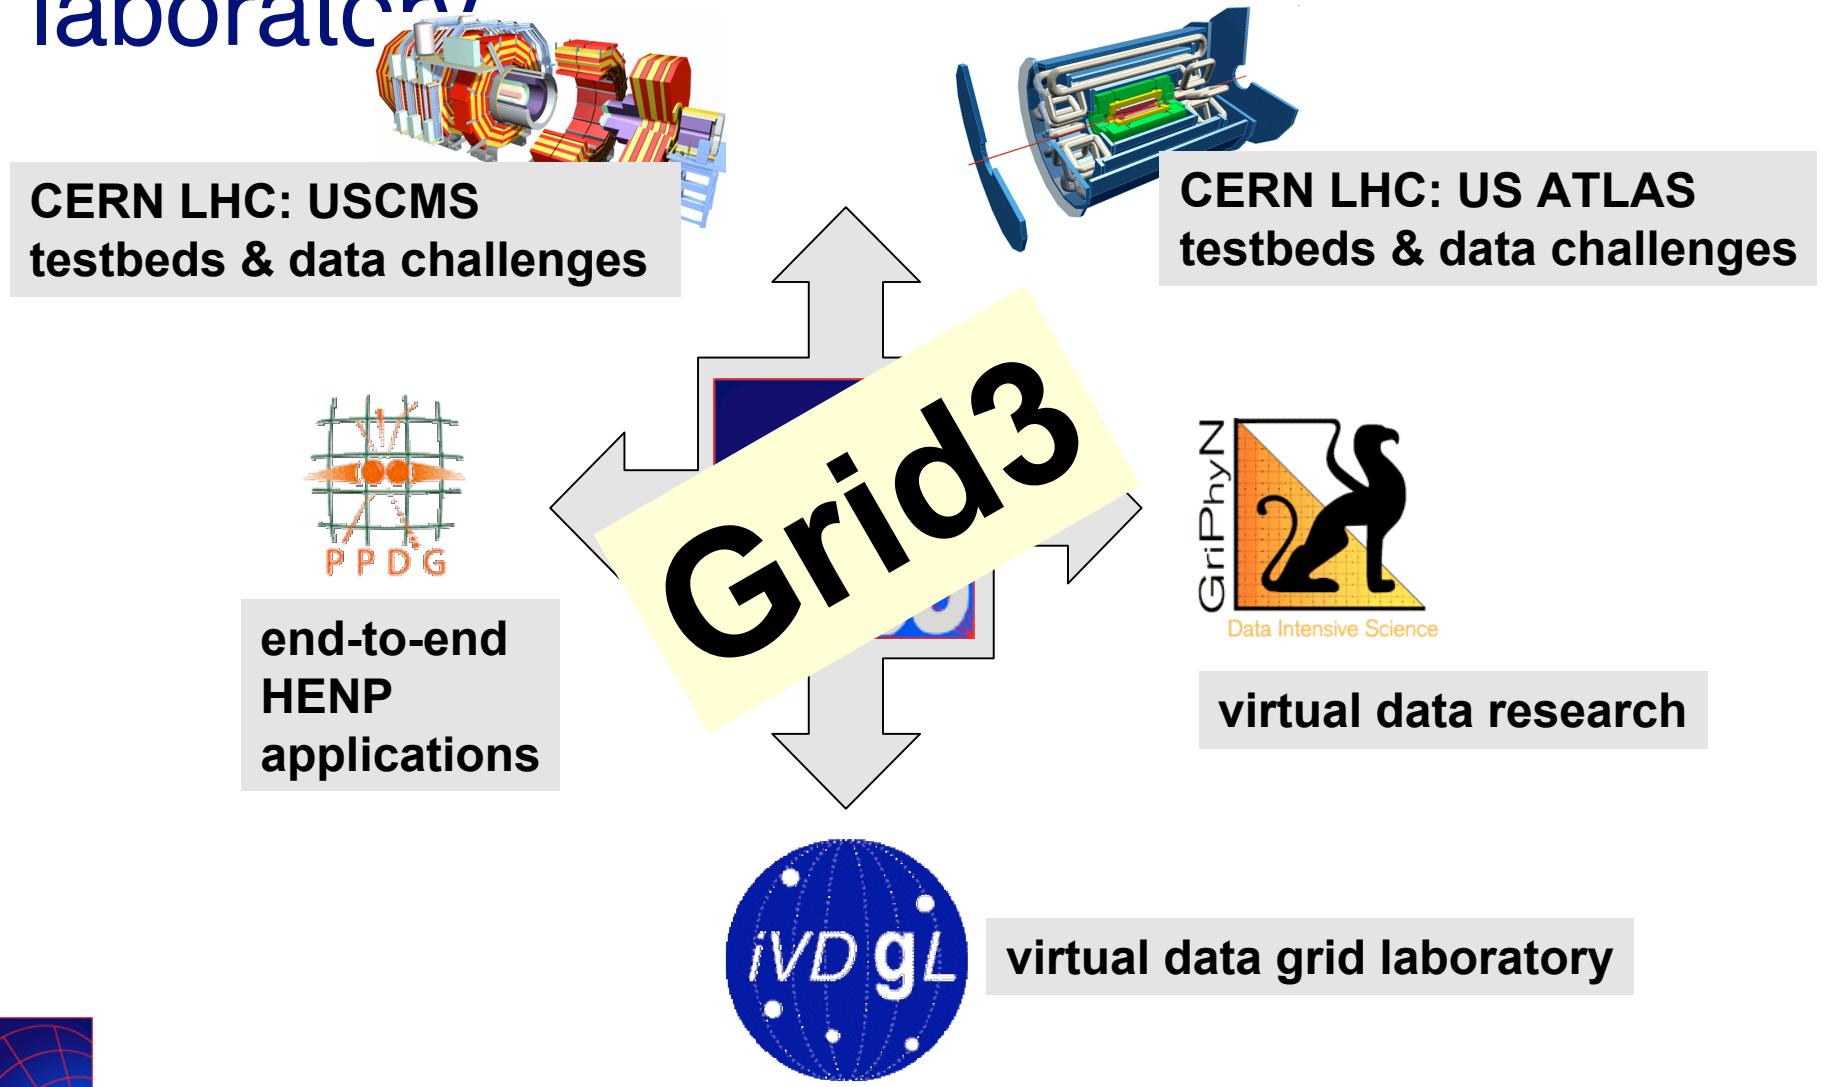
and a few thoughts on federating grids



Rob Gardner
University of Chicago
rwg@hep.uchicago.edu

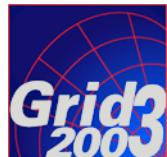
GGF PGM
May 7, 2004

Grid2003: an application grid laboratory



Grid3 at a Glance

- Grid environment built from core Globus and Condor middleware, as delivered through the Virtual Data Toolkit (VDT)
 - GRAM, GridFTP, MDS, RLS, VDS, VOMS, ...
- ...equipped with VO and multi-VO security, monitoring, and operations services
- ...allowing federation with other Grids where possible, eg. CERN LHC Computing Grid (LCG)
 - USATLAS: GriPhyN VDS execution on LCG sites
 - USCMS: storage element interoperability (SRM/dCache)



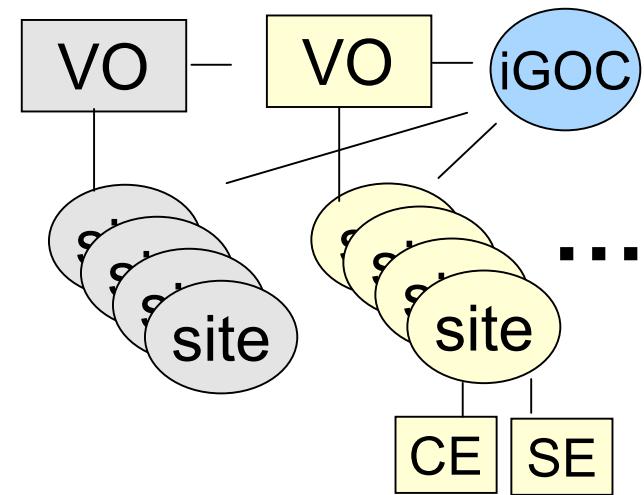
Grid3 – a snapshot of sites



- 27 sites, multi-VO
- shared resources
- ~2000 CPUs
- dynamic – roll in/out

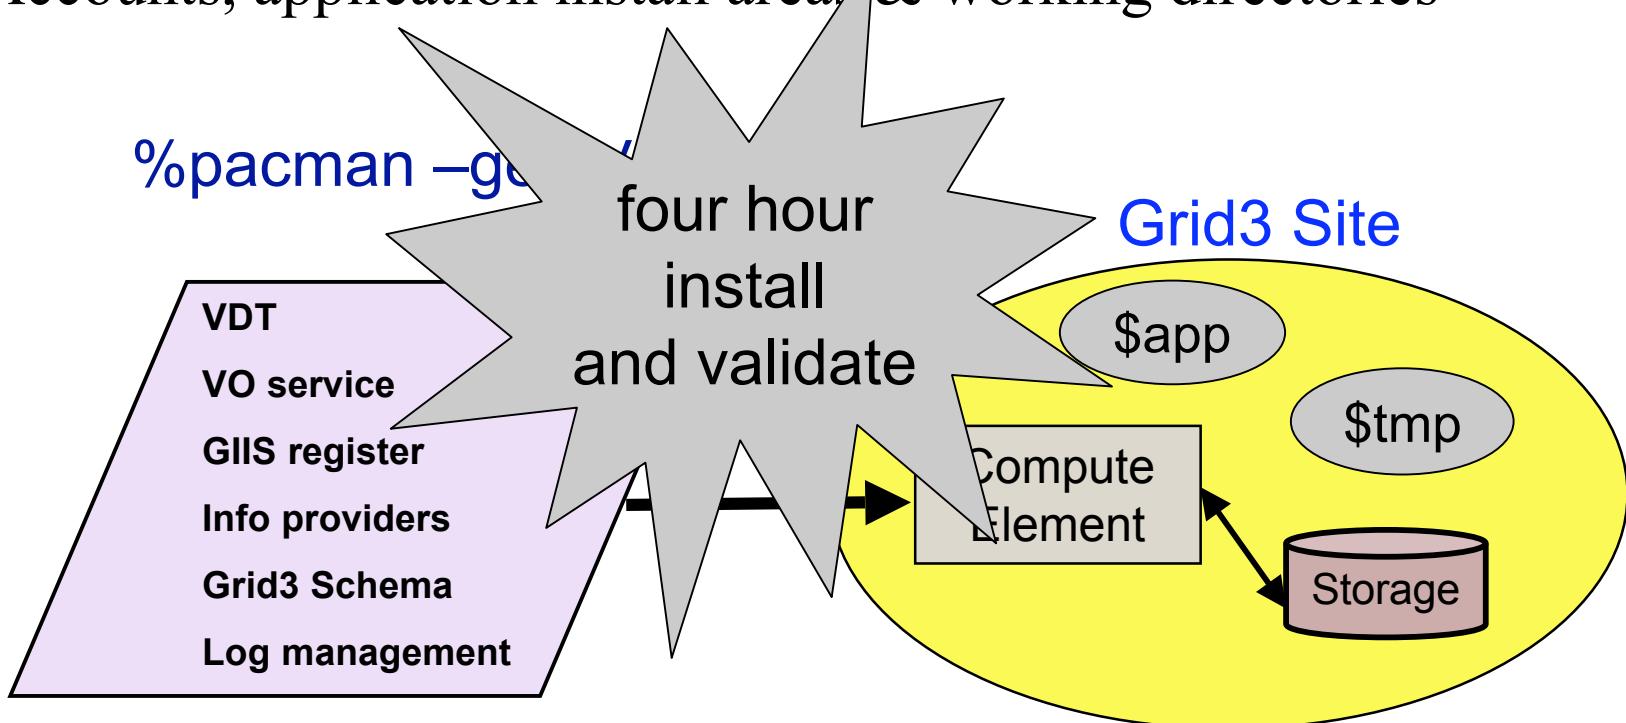
Grid Design

- Simple approach:
 - Sites consisting of
 - Computing element (CE)
 - Storage element (SE)
 - Information and monitoring services
 - VO level, and multi-VO
 - VO information services
 - Operations (iGOC)
 - Minimal use of grid-wide systems
 - No centralized resource broker, replica/metadata catalogs, or command line interface
 - to be provided by individual VO's



Site Services and Installation

- Goal is to install and configure with minimal human intervention
- Use Pacman tool and distributed software “caches”
- Registers site with VO and Grid3 level services
- Accounts, application install areas & working directories



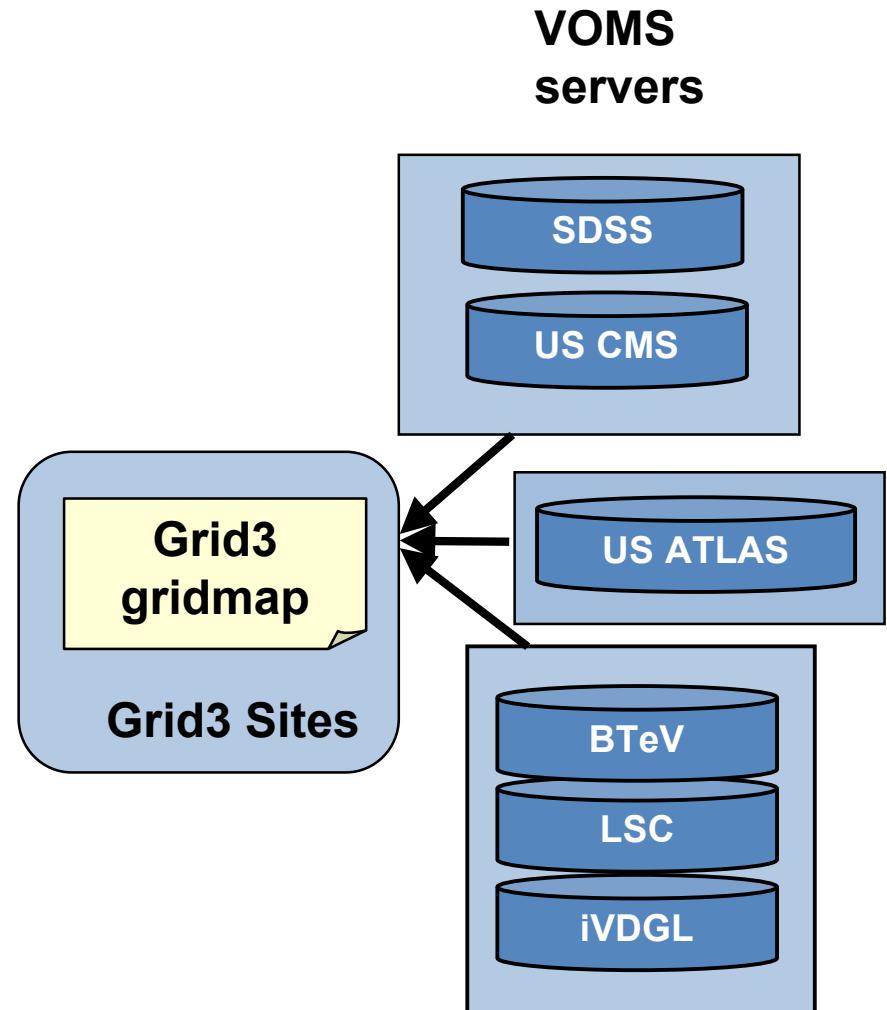
VDT 1.1.4

- Virtual Data System 1.2.3
- ClassAds 0.9.5
- Condor 6.6.1
- EDG CRL Update 1.2.5
- EDG Make Gridmap 2.1.0
- Fault Tolerant Shell (ftsh) 2.0.0
- Globus 2.4.3 plus patches
- GLUE Information providers,
(CVS version 1.79, 25-July-
2003)
- GLUE Schema 1.1, extended
version 1
- GPT 3.1
- GSI-Enabled OpenSSH 3.0
- Java SDK 1.4.1
- KX509 2031111
- Monalisa 0.95
- MyProxy 1.11
- Netlogger 2.2
- PyGlobus 1.0
- PyGlobus URL Copy 1.1.2.11
- RLS 2.1.4
- UberFTP 1.3



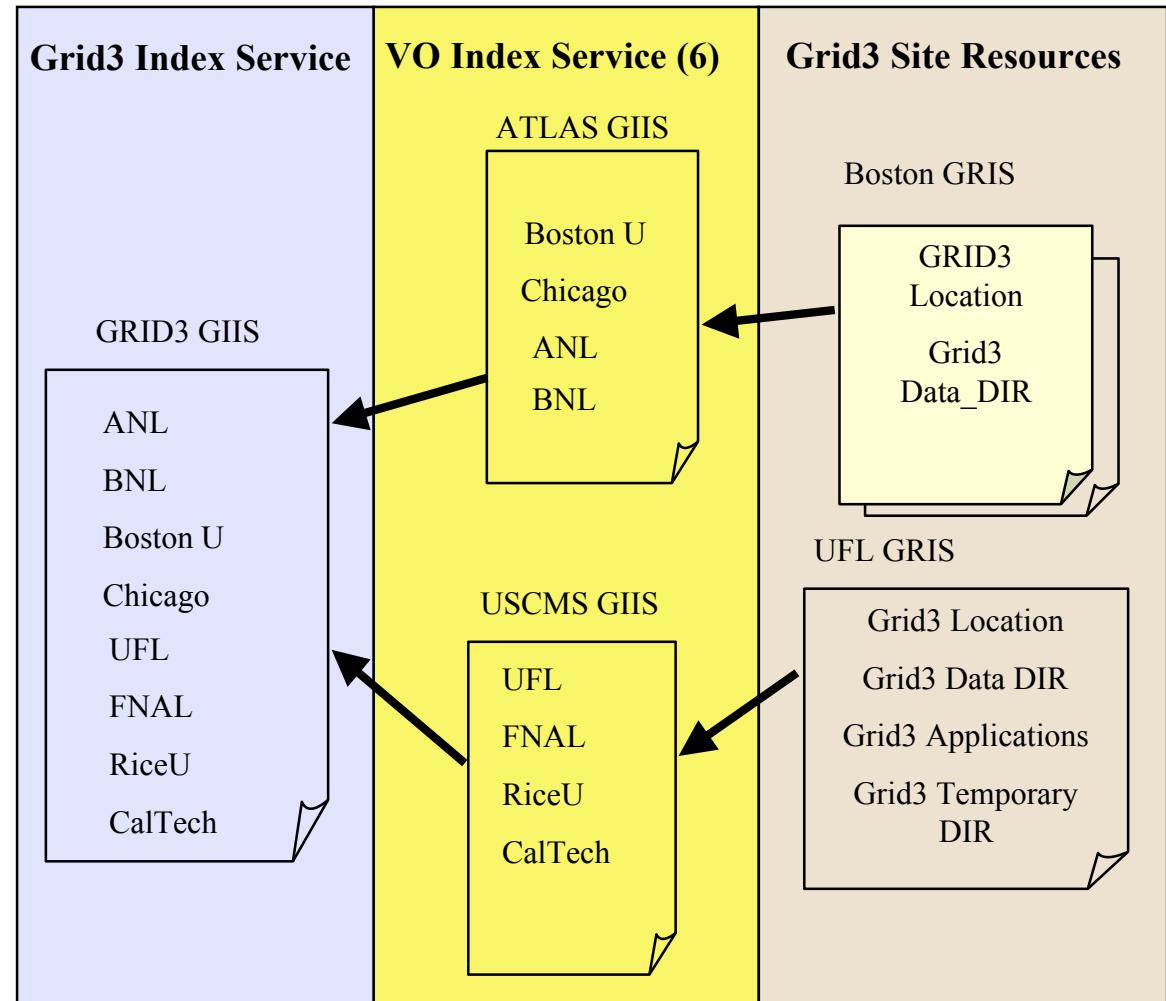
Multi-VO Security Model

- DOEGrids Certificate Authority
- PPDG and iVDGL Registration Authorities, with VO or site sponsorship
- Automated multi-VO authorization, using VOMS
- Each VO manages a service and it's members
- Each Grid3 site is able to generate a Globus gridmap file with an authenticated SOAP query to each VO service
- Site-specific adjustments or mappings



Information Services

- Information Index Service (GIIS)
- Automated Resource registration to Index Service
- MDS Grid3 Schema development & Information Provider verification
- MDS tuning

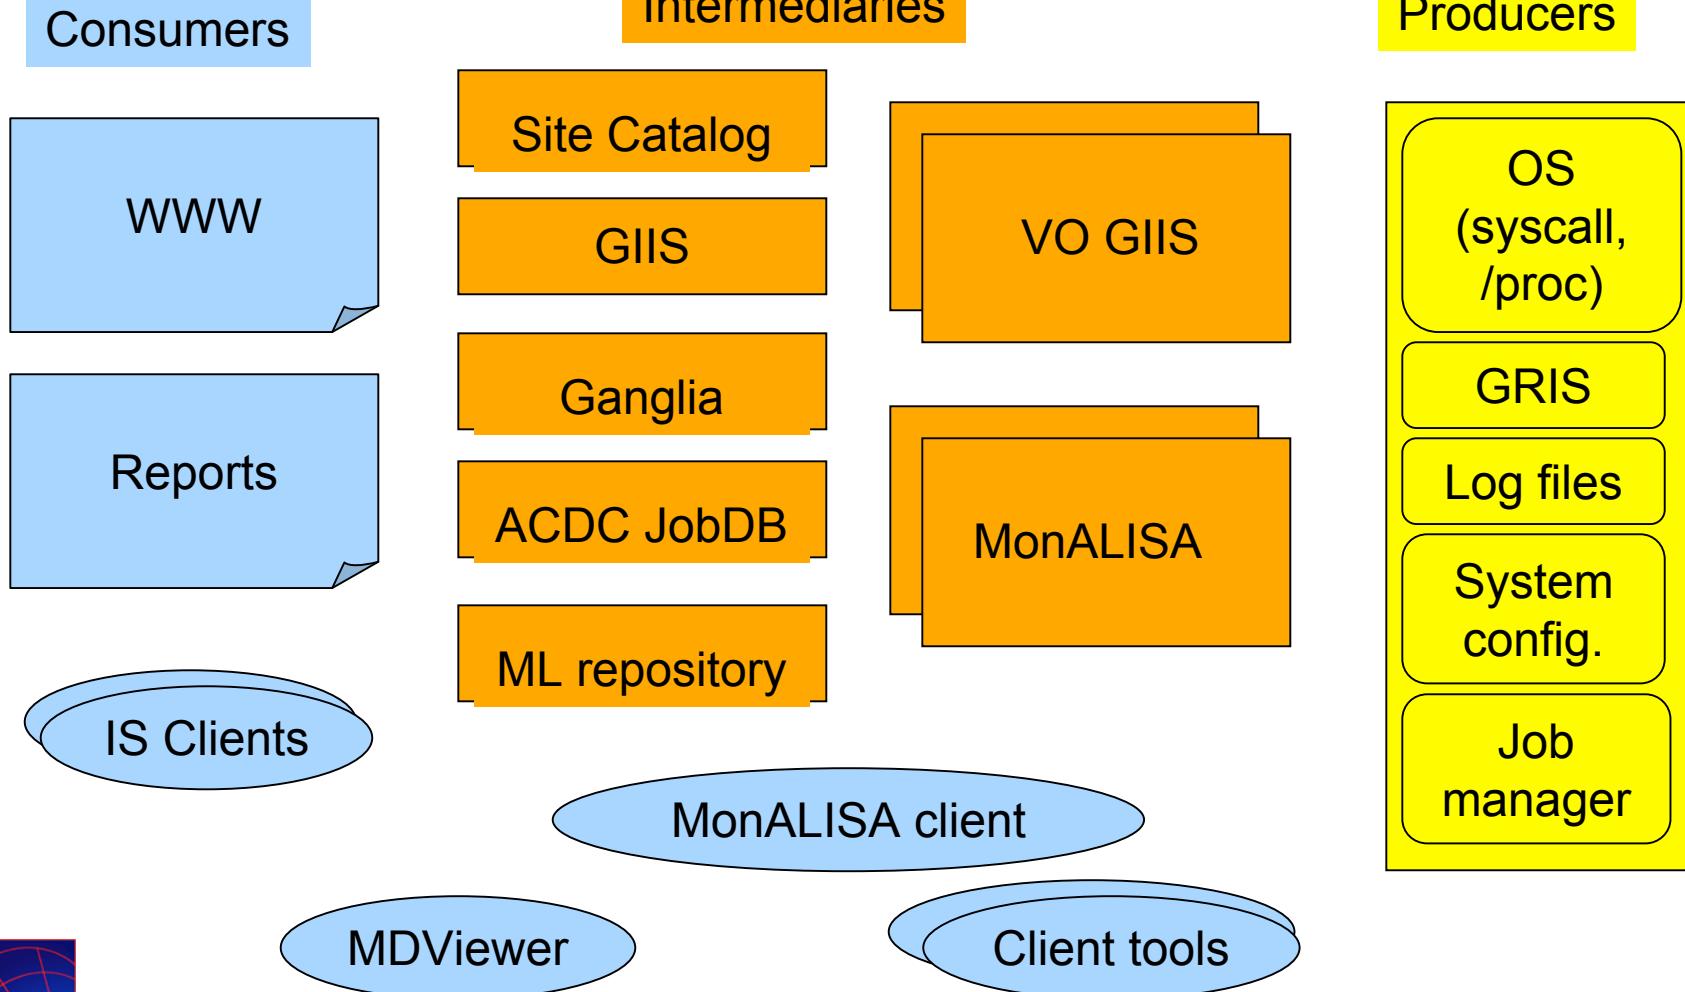


Monitoring Services and Tools

- Globus MDS
 - GLUE and Grid3 schema based information services
- Ganglia
 - cluster monitoring information such as CPU and network load, memory and disk usage
- MonALISA
 - Agent based monitoring gathering system
- ACDC Job Monitoring System
 - Grid submitted jobs to query the job managers and collect information
- Metrics Data Viewer (MDViewer)
 - Plots information collected by the different monitoring systems

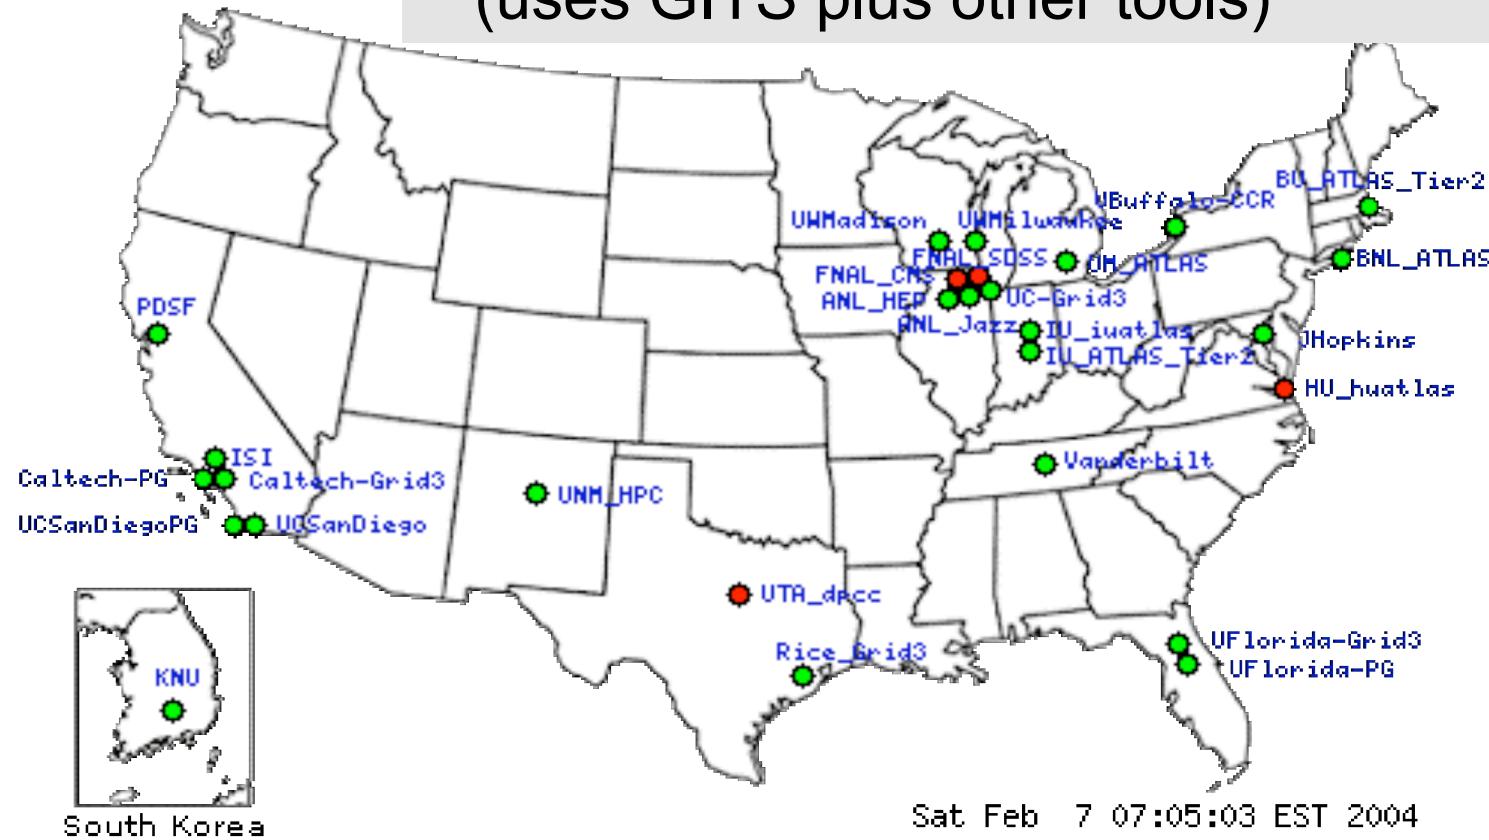


Collective Monitoring Framework



Site Catalog Tool

periodic validation checks & status display
(uses GITS plus other tools)



iVDGL Operations Center (iGOC)

- Co-located with Abilene NOC, hosted by Indiana University
- Hosts/manages multi-VO services:
 - top level Ganglia collectors
 - MonALISA web server and archival service
 - top level GIIS
 - VOMS servers for iVDGL, BTeV, SDSS
 - Site Catalog service
 - iVDGL Pacman caches
- Trouble ticket systems
 - phone (24 hr), web and email based collection and reporting system
 - Investigation and resolution of grid middleware problems at the level of 16-20 contacts per week
- Weekly operations meetings for troubleshooting



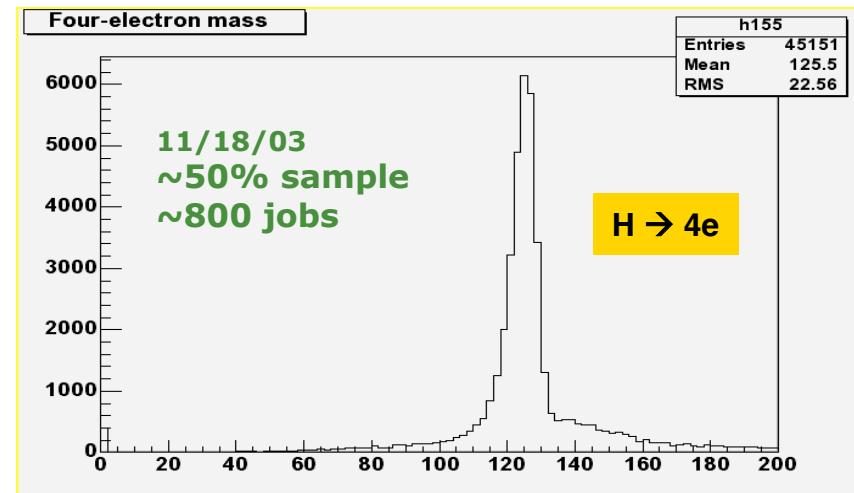
Scientific Applications

- Rule: install through grid jobs, encourage self-contained environments
- 11 applications
 - LHC (ATLAS, CMS)
 - Astrophysics (LIGO/SDSS)
 - Biochemical
 - Molecular X-ray diffraction
 - GADU/Gnare: compares protein sequences
 - Computer Science
 - Adaptive data placement and scheduling algorithms
 - Grid Exerciser
 - Over 100 users authorized to run on Grid3



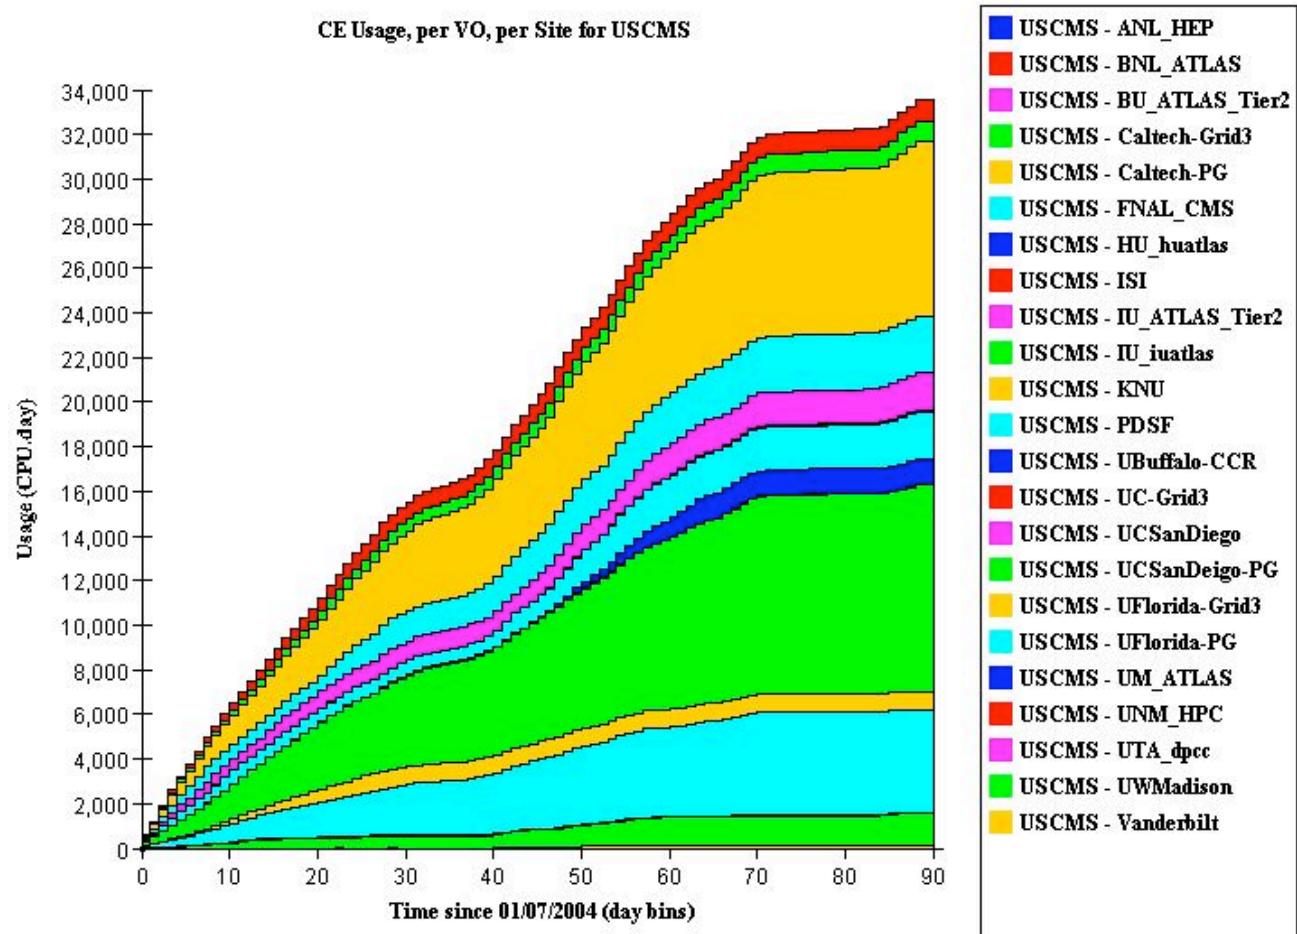
Physics Workflow on Grid3

- Higgs \rightarrow 4 lepton sample
 - Simulate and reconstruct
 - 2000 jobs (X 6 subjobs);
100~200 events per job
- Workflow
 - Jobs defined, stored and registered in RLS
 - Chimera, Pegasus, and Condor-G to submit to Grid3
 - Data pushed to SE at Brookhaven, registered RLS
 - Datasets assembled, passed to physicist for analysis
- Grid3 resources used
 - 16 sites, ~1500 CPUs exercised; peak ~800 jobs over three week period



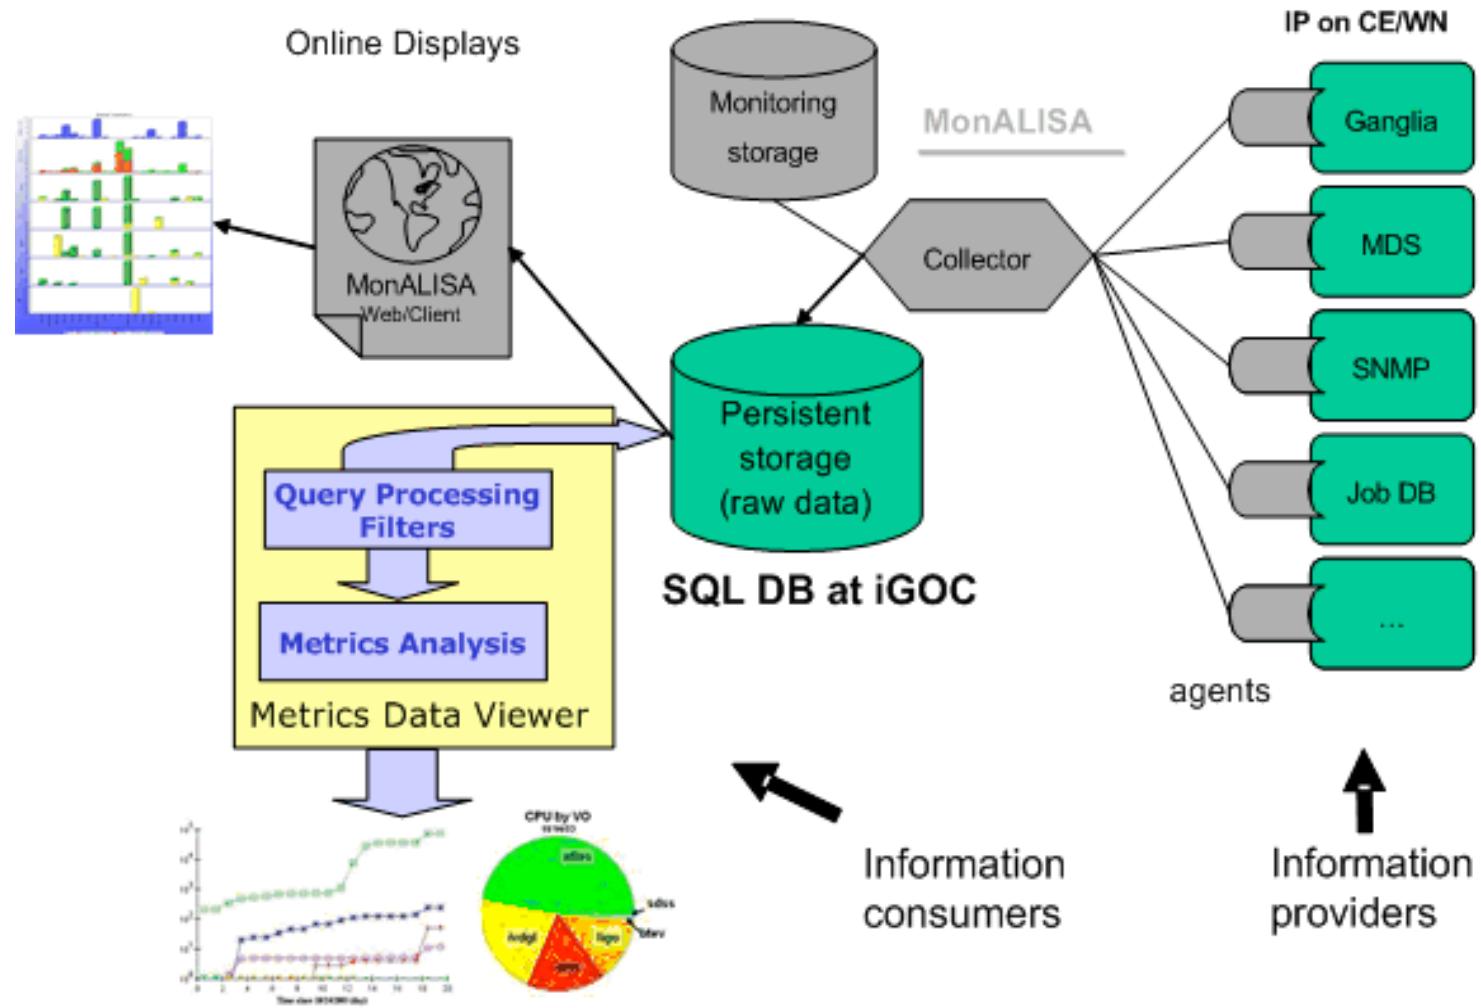
Large Scale Grid3 Operations

- USCMS DC04
- First three months of 2004
- ~40% of non-CMS resources

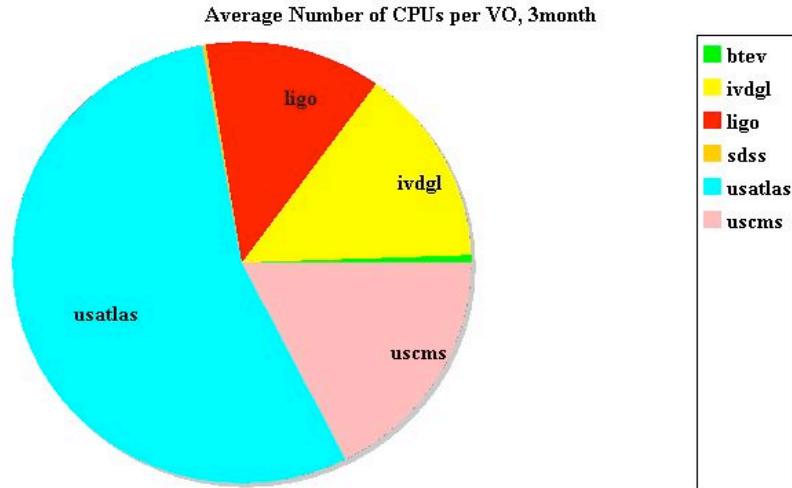


Grid3 Metrics (1)

- Goal: associate activity with VOs and sites



Grid3 Metrics (2)

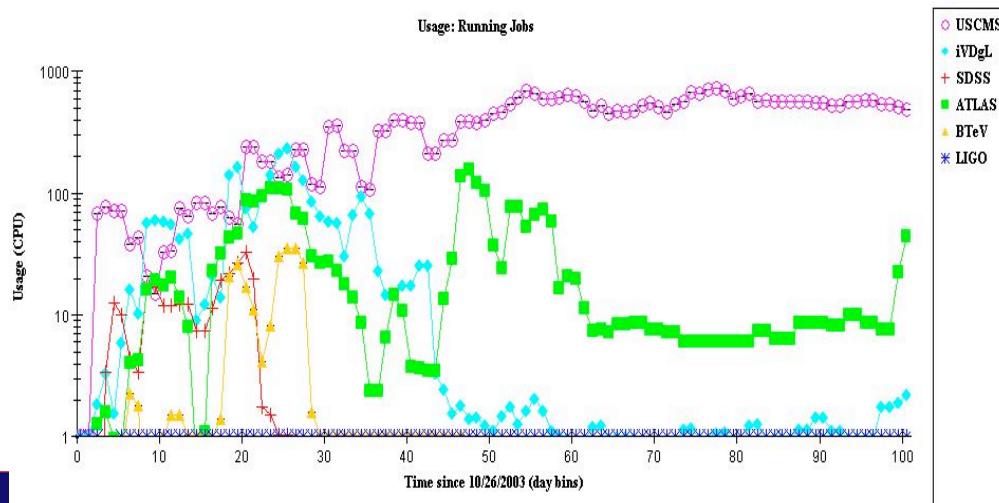


Hardware resources

- ❑ Total of 2762 CPUs
 - Maximum CPU count
 - Off project > 60%
- ❑ Total of 27 sites
 - 27 administrative domains
 - US and Korea

Running jobs

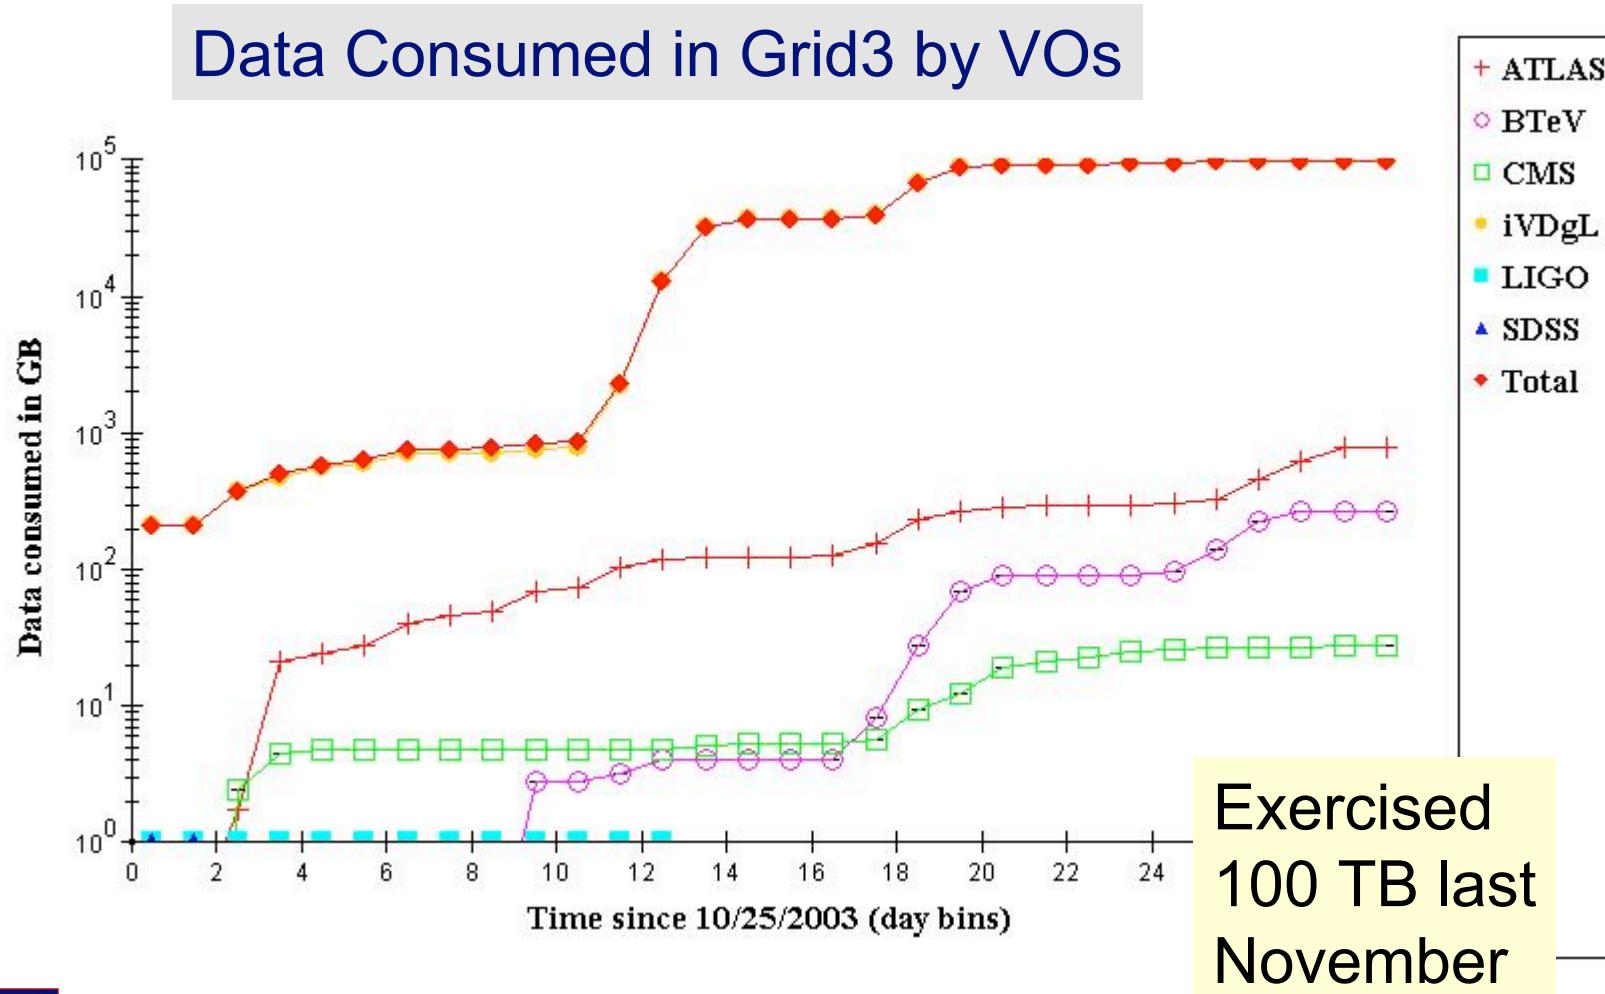
- ❑ Peak number of jobs 1100
- ❑ During SC2003 various applications were running simultaneously across various Grid3 sites



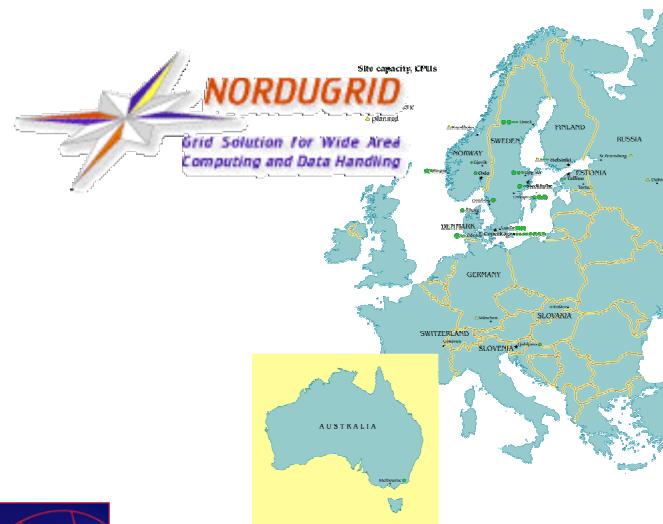
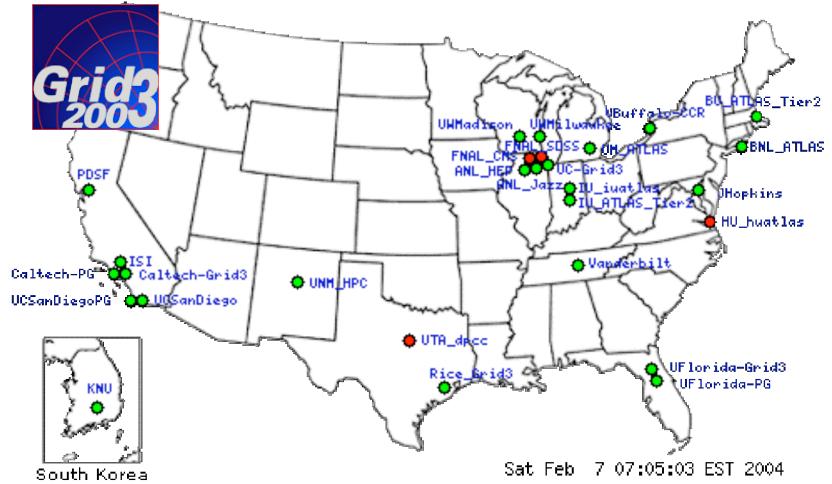
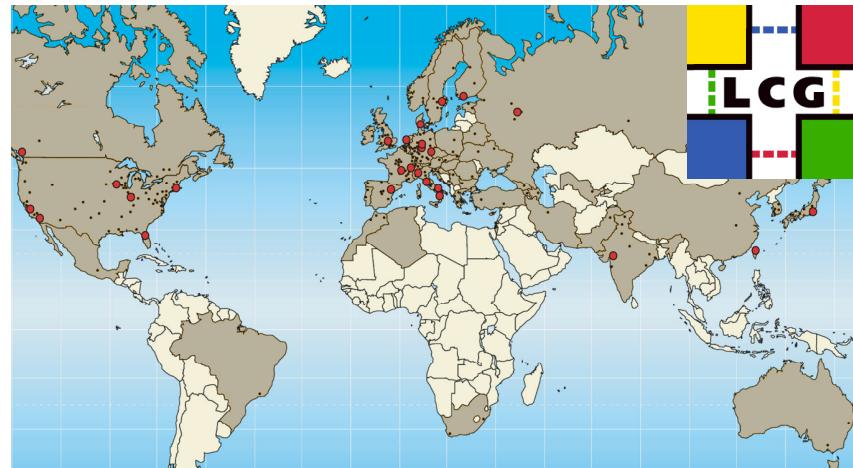
Grid3 Metrics (3)



Grid3 Metrics (4)



Federation of Grids



Few thoughts on Federation

- First and foremost, scientists belong to VOs which span Grids...
 - ...an LHC scientist may wish (and be authorized to) access resources from the LHC Computing Grid (LCG), UK Grid, NorduGrid, APGrid, participating US Grids like Grid3 and its successors (i.e. Open Science Grid)
 - US LHC physicists may wish to gain opportunistic access to other resources like TeraGrid, campus-based Grids, etc
- What is the best way to enable this? Many issues.
 - authorization, accounting, catalogs, (workflow)

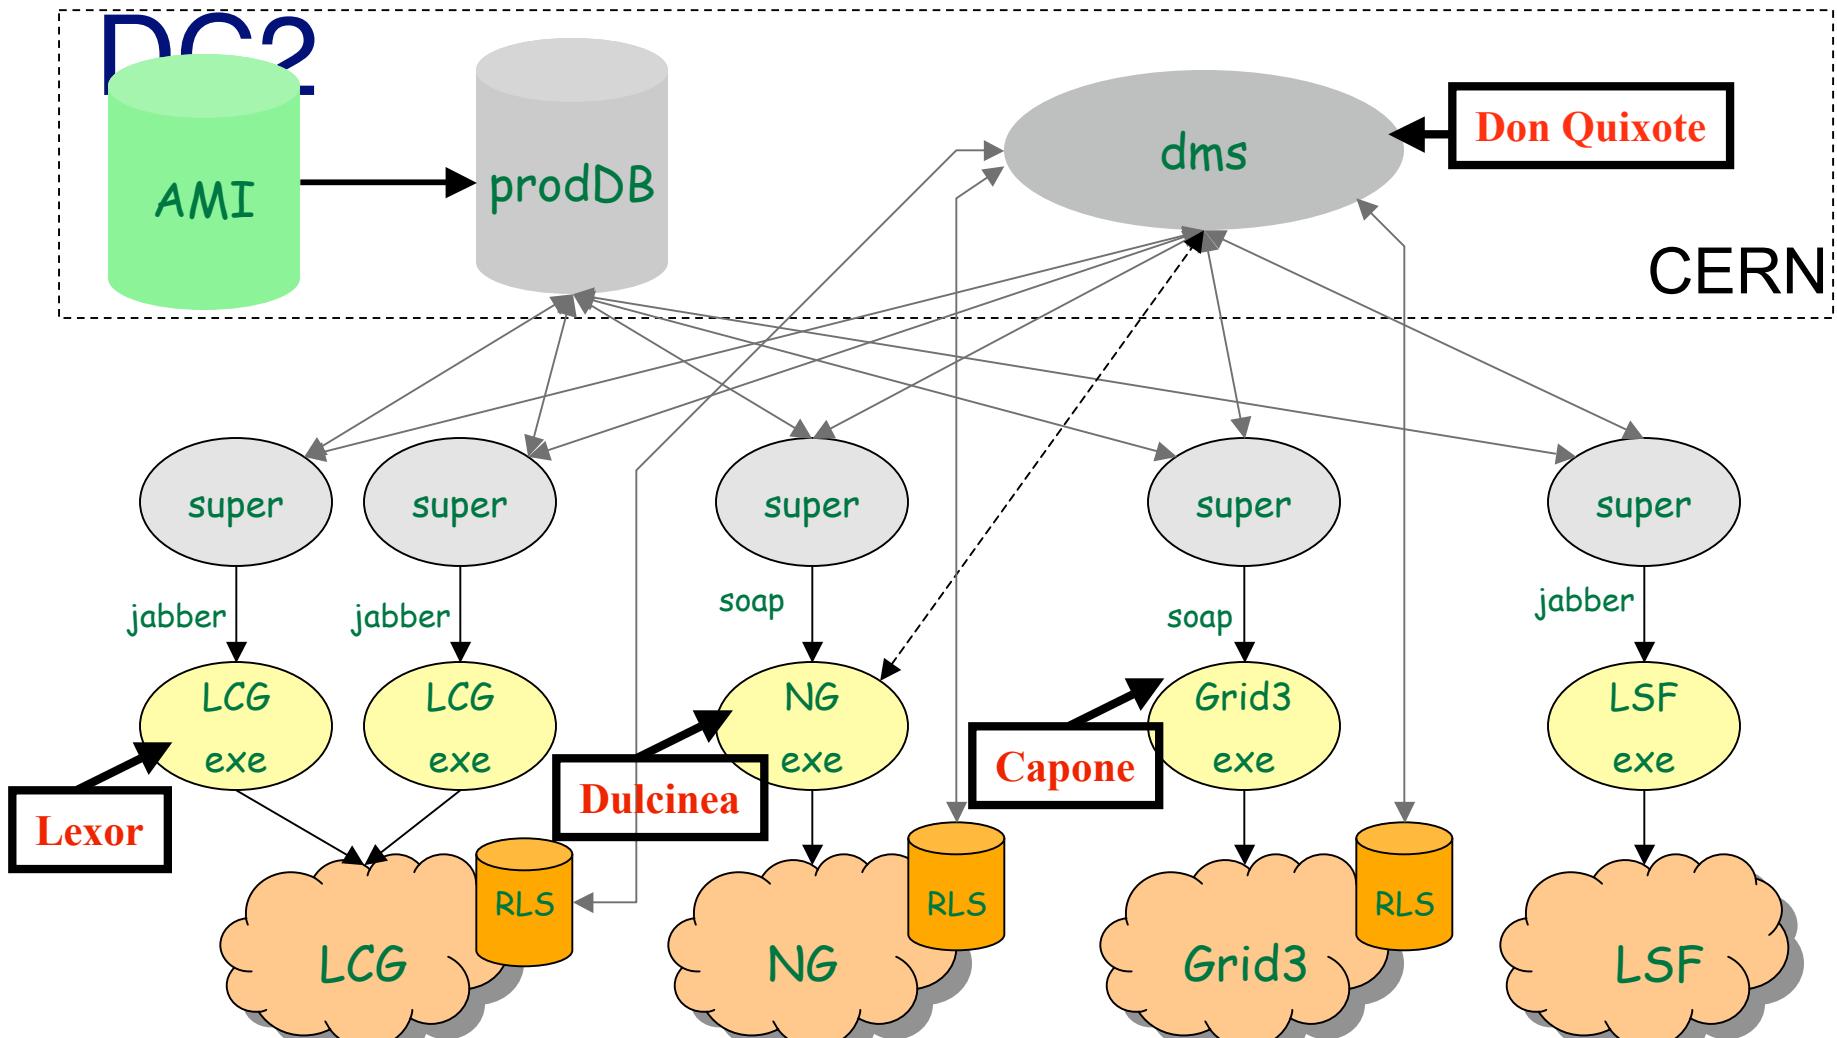


Workflow between Grids



- Workflow exercise Grid3 and LCG
 - Grid3 produced output staged on disk at Brookhaven lab in NY
 - Use VDT client to submit to an LCG server (bypass RB)
 - Jobs executing on LCG:
 - Input files registered at BNL-RLI
 - Stage data from BNL to local scratch area on LCG grid
 - Run Athena-reconstruction on LCG
 - Write data to local storage element on LCG
 - Copy back to disk cache at NY, RLS register
 - 3rd-party transfer between LCG SE and BNL gridftp servers
 - Post job: Dataset catalog and analysis
- Many lessons learned about site architectural differences, service configurations; *Concrete steps towards combined use of LCG and US Grid3*

ATLAS Production System for



system implemented, production starting this week ⁴

Conclusions

- Grid2003 project created a persistent, production scale grid infrastructure, Grid3
- Continuous use since 10/03
- Grid3 will undergo adiabatic upgrades, and continue operations for 2004 (Grid3+)
- Outlook:
 - Follow-on project will involve more sites and VOs
 - Larger “consortia” forming, Open Science Grid (OSG)
 - New projects will be defined in this larger framework



Acknowledgements

- the entire Grid2003 team which did all the work!
 - site administrators
 - application developers
 - iGOC team
 - coordinators and project teams
- Slides I've taken from presentations by:
 - Leigh Grundhoefer (Indiana, iVDGL)
 - Ian Fisk (Fermilab, USCMS)
 - Ruth Pordes (Fermilab, PPDG)
 - Jorge Rodriguez (Florida, USCMS/iVDGL)

